

DA5030.Proj.Sui

Xin (Sue) Sui

2020-04-15

CRISP-DM Business Understanding

Bank churning is defined as movement of customers from one bank to another due to several reasons such as: low interest rates and fees, customer service, latest technology, store hours and locations. As we know, it is more costly to sign in new clients than retaining an existing one. By the time customers churn, banks lose money and it is too late to know the reason.

The goal is to build a model to predict whether a customer will churn or not based on customers' background and financial status.

- Predict whether a customer will churn or not.
- Help banks determine and develop a churn prevention plan to prevent customers from churning.
- Reduce losses to banks.
- Better serve customers and retain them.

CRISP-DM Data Understanding

1. Data Acquisition:

- acquisition of data (e.g., CSV or flat file)

This is the data set used in the section "ANN (Artificial Neural Networks)" of the Udemy course from Kirill Eremlenko (Data Scientist & Forex Systems Expert) and Hadelin de Ponteves (Data Scientist), called Deep Learning A-Z™: Hands-On Artificial Neural Networks.

This data was obtained from Kaggle: <https://www.kaggle.com/adammaus/predicting-churn-for-bank-customers>
(<https://www.kaggle.com/adammaus/predicting-churn-for-bank-customers>)

```
# 1. Load the data into dataframe
df <- read.csv("Churn_Modelling.csv", header = TRUE, stringsAsFactors = TRUE)
```

2. Data Exploration:

- Exploratory data plots

- Detection of outliers

- Correlation/collinearity analysis

```
# 1. Check data
head(df)
```

```
str(df)
```

```
## 'data.frame': 10000 obs. of 14 variables:
## $ RowNumber : int 1 2 3 4 5 6 7 8 9 10 ...
## $ CustomerId : int 15634602 15647311 15619304 15701354 15737888 15574012 15592531 15656148 15792365 1559
2389 ...
## $ Surname : Factor w/ 2932 levels "Abazu","Abbie",...: 1116 1178 2041 290 1823 538 178 2001 1147 1082
...
## $ CreditScore : int 619 608 502 699 850 645 822 376 501 684 ...
## $ Geography : Factor w/ 3 levels "France","Germany",...: 1 3 1 1 3 3 1 2 1 1 ...
## $ Gender : Factor w/ 2 levels "Female","Male": 1 1 1 1 1 2 2 1 2 2 ...
## $ Age : int 42 41 42 39 43 44 50 29 44 27 ...
## $ Tenure : int 2 1 8 1 2 8 7 4 4 2 ...
## $ Balance : num 0 83808 159661 0 125511 ...
## $ NumOfProducts : int 1 1 3 2 1 2 2 4 2 1 ...
## $ HasCrCard : int 1 0 1 0 1 1 1 1 0 1 ...
## $ IsActiveMember : int 1 1 0 0 1 0 1 0 1 1 ...
## $ EstimatedSalary: num 101349 112543 113932 93827 79084 ...
## $ Exited : int 1 0 1 0 0 1 0 1 0 0 ...
```

```
summary(df)
```

```
##      RowNumber      CustomerId      Surname      CreditScore
## Min.   :    1      Min.   :15565701      Smith   :   32      Min.   :350.0
## 1st Qu.: 2501      1st Qu.:15628528      Martin   :   29      1st Qu.:584.0
## Median : 5000      Median :15690738      Scott    :   29      Median :652.0
## Mean   : 5000      Mean   :15690941      Walker    :   28      Mean   :650.5
## 3rd Qu.: 7500      3rd Qu.:15753234      Brown     :   26      3rd Qu.:718.0
## Max.   :10000      Max.   :15815690      Genovese  :   25      Max.   :850.0
##                                     (Other) :9831
##      Geography      Gender      Age      Tenure      Balance
## France :5014      Female:4543      Min.    :18.00      Min.    : 0.000      Min.    :    0
## Germany:2509      Male  :5457      1st Qu.:32.00      1st Qu.: 3.000      1st Qu.:    0
## Spain   :2477                                     Median :37.00      Median : 5.000      Median : 97199
##                                     Mean    :38.92      Mean    : 5.013      Mean    : 76486
##                                     3rd Qu.:44.00      3rd Qu.: 7.000      3rd Qu.:127644
##                                     Max.    :92.00      Max.    :10.000      Max.    :250898
##
##      NumOfProducts      HasCrCard      IsActiveMember      EstimatedSalary
## Min.    :1.00      Min.    :0.0000      Min.    :0.0000      Min.    :   11.58
## 1st Qu.:1.00      1st Qu.:0.0000      1st Qu.:0.0000      1st Qu.: 51002.11
## Median :1.00      Median :1.0000      Median :1.0000      Median :100193.91
## Mean    :1.53      Mean    :0.7055      Mean    :0.5151      Mean    :100090.24
## 3rd Qu.:2.00      3rd Qu.:1.0000      3rd Qu.:1.0000      3rd Qu.:149388.25
## Max.    :4.00      Max.    :1.0000      Max.    :1.0000      Max.    :199992.48
##
##      Exited
## Min.    :0.0000
## 1st Qu.:0.0000
## Median :0.0000
## Mean    :0.2037
## 3rd Qu.:0.0000
## Max.    :1.0000
##
```

```
table(df$Exited)
```

```
##
##      0      1
## 7963 2037
```

```
# Has 10000 observations and 14 variables (11 are usable)

# Note: RowNumber, CustomerID and Surname are not useable, and will be eliminated
# Numerical features: CreditScore, Age, Tenure, Balance, NumOfProducts, EstimatedSalary
# Categorical features: Geography, Gender, HasCrCard, IsActiveMember and Exited

numeric <- c("CreditScore", "Age", "Tenure", "Balance", "NumOfProducts", "EstimatedSalary")
categorical <- c("Geography", "Gender", "HasCrCard", "IsActiveMember", "Exited")

# 2. Check for missing and NA values, there are none
sum(is.na(df))
```

```
## [1] 0
```

```
length(which(df == "?"))
```

```
## [1] 0
```

```
length(which(df == "NA"))
```

```
## [1] 0
```

```
length(which(df == "N/A"))
```

```
## [1] 0
```

```
# 3. Check for outliers
# Create z-score standarization function.
z_normalize <- function(x) {
  return ((x - mean(x)) / sd(x))
}

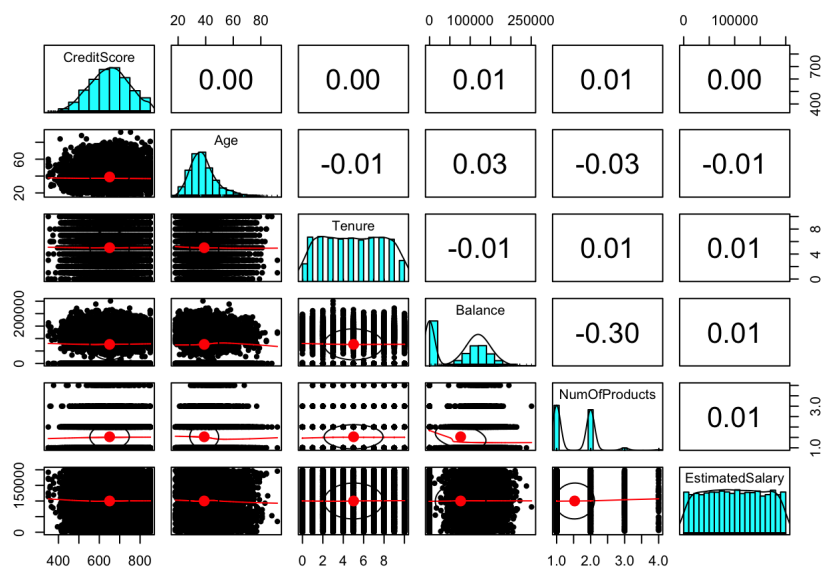
# Normalize the numerical features.
numeric <- df[c(4,7:10,13)]
norm <- apply(numeric, 2, z_normalize)

# Find the outliers (outliers are considered 3 standard deviations away from the mean).
outliers <- abs(norm) > 3
sum(outliers) # 201 outliers
```

```
## [1] 201
```

```
outliers_column <- which(apply(outliers, 1, function(x) sum(x)!=0))
```

```
# 4. Check for correlation and collinearity:
pairs.panels(numeric)
```



Comment: In the my dataset, there are no missing values. If there were missing data and not much, I would remove them and state it. If there were a lot of missing data, I would impute them with either a value between min and max, but this might cause high variance and poor fit. Or impute with average, or median of similar data by clustering. There are 201 outliers which I will remove which might increase variability when training the models. The algorithms I plan to use are Naive Bayes, Decision Trees, Neural Network and SVM, removing the outliers will help improve model performance overall. Since none of the four models I will be building are statistical learners, following a Gaussian distribution for the data is not required.

CRISP-DM Data Understanding and Data Preparation

1. Data Cleaning & Shaping:

- Data imputation
- Normalization/standardization of feature values
- Feature engineering: dummy codes
- Feature engineering: PCA
- Feature engineering: new derived features

```
# Data cleaning

# 1. Check for useful features: First three columns are unique row number,
# customer id and their surname, I will exclude from the data.
df <- df[c(-1:-3)]

# 2. Convert categorical features to factors
df$Exited <- ifelse(df$Exited == 0, "no", "yes")
df[categorical] <- lapply(df[categorical], factor)
head(df)
```

```

# 3. Data imputation: Remove outliers
df <- df[c(-outliers_column), ]

# Data shaping: Here I will shape the data based on each model I use, some
# require normalization of numerical features; some require conversion to
# categorical features; some require conversion to numerical features.

# 4. Normalization/standardization of feature values:

# 4.1 Normalization: Normalize for Neural Network (min-max) it works best when
# input data are scaled.
numeric <- c("CreditScore", "Age", "Tenure", "Balance", "NumOfProducts", "EstimatedSalary")
normalize <- function(x) {
  return((x - min(x))/(max(x) - min(x)))
}
nn_norm <- as.data.frame(lapply(df[numeric], normalize))

nn_df <- cbind(nn_norm, y = df$Exited)

# 4.3 Normalization: Normalize for SVM (min-max): it works best when input data
# are scaled.
normalize <- function(x) {
  return((x - min(x))/(max(x) - min(x)))
}
svm_norm <- as.data.frame(lapply(df[numeric], normalize))

# 5. Feature engineering:

# 5.1 Dummy code categorical features to numeric for SVM
dum <- df[categorical] %>% select(-Exited)
dmy <- dummyVars("~.", dum)
svm_dmy <- data.frame(predict(dmy, newdata = dum))

svm_final_df <- cbind(svm_norm, svm_dmy, y = df$Exited)


# 5.2 New derived features for Naive Bayes by binning as it uses only categorical
# features
nb_df <- df
nb_df$CreditScore <- bin(nb_df$CreditScore, nbins = 5, labels = c("1", "2", "3",
  "4", "5"))
nb_df$Age <- bin(nb_df$Age, nbins = 5, labels = c("1", "2", "3", "4", "5"))
nb_df$Tenure <- bin(nb_df$Tenure, nbins = 5, labels = c("1", "2", "3", "4", "5"))
nb_df$Balance <- bin(nb_df$Balance, nbins = 2, labels = c("1", "2"))
nb_df$NumOfProducts <- bin(nb_df$Tenure, nbins = 4, labels = c("1", "2", "3", "4"))
nb_df$EstimatedSalary <- bin(nb_df$EstimatedSalary, nbins = 5, labels = c("1", "2",
  "3", "4", "5"))
head(nb_df)

# 5.3. Feature engineering: PCA
cor_df <- df
pca_num_df <- cor_df[c("CreditScore", "Age", "Tenure", "Balance", "NumOfProducts",
  "EstimatedSalary")]
pca_num <- apply(pca_num_df, 2, function(x) as.numeric(as.character(x))) # Convert to all numbers

cor_num <- cor(pca_num)

ggplot(data = melt(cor_num), aes(Var1, Var2, fill = value)) + geom_tile(colour = "white") +
  scale_fill_viridis_c(name = "correlation index") + theme(axis.text.x = element_text(angle = 90,
    hjust = 1)) + scale_x_discrete(labels = abbreviate) # Age

```



```
pca_numeric <- prcomp(pca_num_df, center = TRUE, scale = TRUE)
summary(pca_numeric)
```

```
## Importance of components:
##              PC1      PC2      PC3      PC4      PC5      PC6
## Standard deviation  1.1595 1.0082 1.0006 0.9959 0.9895 0.8168
## Proportion of Variance 0.2241 0.1694 0.1668 0.1653 0.1632 0.1112
## Cumulative Proportion 0.2241 0.3935 0.5603 0.7256 0.8888 1.0000
```

```
# In this case, I did PCA on numerical features and found that Credit Score and
# Age are the important two factors, they explain about 40% of the total
# variation in the data.
```

CRISP-DM Data Modeling

In This section, holdout method is used to construct four different models, and each model is tuned to see whether or not it improves performance. Models built are then used to predict test data set and accuracy is calculated.

Model Construction:

1. Creation of training & validation subsets
2. Construction of at least three related models:

```
2.1 - Neural Network (Parametric)
2.2 - Decision Trees (Non-parametric)
2.3 - Support Vector Machine (Parametric)
2.4 - Naive Bayes (Parametric)
```

2.1 - I chose Neural Network because it can be used for both regression and classification.

2.2 - I chose SVM because it is good for binary classification, it attempts to find a hyper-plane separating the different classes of the training instances, with the maximum error margin.

2.3 - I choose Decision Trees because it doesn't require transformation of data and accepts both categorical and numerical features, it is very simple, fast and efficient.

2.4 - I choose Naive Bayes as a comparison to the previous three and see how binning the numerical features perform in the overall model.

1. Creation of training & validation subsets

```
# set.seed for reproducibility.
set.seed(123)

# Hold out method using stratified holdout sampling:
split <- createDataPartition(df$Exited, p = 0.75, list = FALSE)

training <- df[split, ]
testing <- df[-split, ]

table(training$Exited) %>% prop.table
```

```
##
##          no          yes
## 0.8001361 0.1998639
```

```
table(testing$Exited) %>% prop.table
```

```
##
##          no          yes
## 0.8003267 0.1996733
```

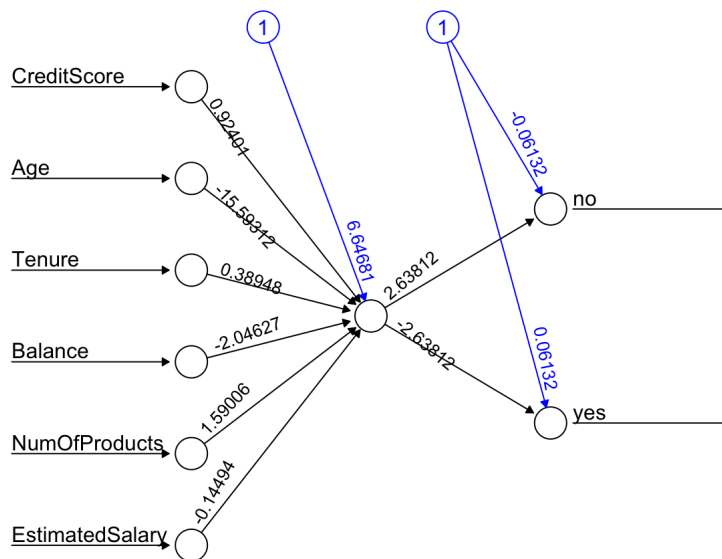
2.1. Construction of Neural Network

```
set.seed(123)

# 1. Split the data (Holdout method)
train <- nn_df[split, ]
test <- nn_df[-split, ]

# 2. Construct Neural Network classifier (for binary classification, act.fct =
# logistic is used here)
nn <- neuralnet(y ~ ., data = train, hidden = 1, linear.output = FALSE, err.fct = "ce",
  act.fct = "logistic", likelihood = TRUE)

plot(nn, rep = "best")
```



Error: 6319.444014 Steps: 2439

```
result <- compute(nn, test[-7])
nn_prediction <- as.factor(ifelse(result$net.result[, 1] < 0.5, "yes", "no"))

confusionMatrix(nn_prediction, test$y, mode = "prec_recall")
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction   no  yes
##           no 1863 395
##           yes  97  94
##
##           Accuracy : 0.7991
##           95% CI : (0.7827, 0.8148)
##           No Information Rate : 0.8003
##           P-Value [Acc > NIR] : 0.5721
##
##           Kappa : 0.1851
##
## Mcnemar's Test P-Value : <2e-16
##
##           Precision : 0.8251
##           Recall : 0.9505
##           F1 : 0.8834
##           Prevalence : 0.8003
##           Detection Rate : 0.7607
##           Detection Prevalence : 0.9220
##           Balanced Accuracy : 0.5714
##
##           'Positive' Class : no
##
```

```
# Accuracy : 0.7991 Precision : 0.8251 Recall : 0.9505 F1 : 0.8834
```


```
# Accuracy: measure of all the correctly identified cases. Precision: measure of
# the correctly identified positive cases from all the predicted positive cases.
# Recall: measure of the correctly identified positive cases from all the actual
# positive cases. F-score: measure of test accuracy. It combines both the
# precision and the recall using the harmonic mean, it describes the model
# performance. It gives a better measure of the incorrectly classified cases than
# the Accuracy Metric.
```

```
# AUC-----
prob_nn <- result$net.result
colAUC(prob_nn, test$y)
```

```
##           [,1]      [,2]
## no vs. yes 0.7485216 0.7485216
```




```
colAUC(prob_nn, test$y, plotROC = TRUE)
```



```
##           [,1]      [,2]
## no vs. yes 0.7485216 0.7485216
```

```
auc_nn <- roc(response = test$y, predictor = result$net.result[, 1])
plot(auc_nn)
```



```
auc_nn$auc # 74.85%
```

```
## Area under the curve: 0.7485
```

```
# Comment: Here I choose to use all numeric features from the data to build
# neural network classifier because the with the the rest of categorical
# features, the computer runs so much and takes up all the CPU and takes a very
# long run time. (I've tried to use both numerical and categorical features, but
# the has a very long run time and uses a lot of computer memory). According to
# the book and Professor's notes, Neural Network can take in both numerical and
# categorical features and automatically converts them to dummy code, and assigns
# bias and weights to each input and back-propagates, it is considered as a black
# box, NP-complete problem which takes an exponential amount of time to run with
# the amount of input data, it takes up all the CPU and very computationally
# expensive. I will demonstrate tuning parameter for Neural Network in the
# following k-fold cross validation step. Accuracy for improved model is 83.54%,
# F-score is 88.34% and AUC is 80.15%.
```

2.2. Construction of Support Vector Machine

```
set.seed(123)

# Hold out method using stratified holdout sampling:
split <- createDataPartition(df$Exited, p = 0.75, list = FALSE)

train <- svm_final_df[split,]
test <- svm_final_df[-split,]

svm_classifier_l <- svm(y ~ ., data = train, kernel="linear", scaled = TRUE, probability = TRUE)

pred_svm_l <- predict(svm_classifier_l, test[-16], decision.values = TRUE, probability = TRUE)

svm_prediction_l <- predict(svm_classifier_l, test[-16], type = "prob")

confusionMatrix(svm_prediction_l, test$y, mode = "prec_recall")
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction   no  yes
##           no 1960 489
##           yes    0   0
##
##           Accuracy : 0.8003
##           95% CI : (0.7839, 0.816)
##           No Information Rate : 0.8003
##           P-Value [Acc > NIR] : 0.5121
##
##           Kappa : 0
##
## Mcnemar's Test P-Value : <2e-16
##
##           Precision : 0.8003
##           Recall : 1.0000
##           F1 : 0.8891
##           Prevalence : 0.8003
##           Detection Rate : 0.8003
##           Detection Prevalence : 1.0000
##           Balanced Accuracy : 0.5000
##
##           'Positive' Class : no
##
```

```
# Accuracy : 0.8003
# Precision : 0.8003
# Recall : 1.0000

# Improvement with kernel = radial
svm_classifier_k <- svm(y ~ ., data = train, kernel="radial", scaled = TRUE, probability = TRUE)

pred_svm_k <- predict(svm_classifier_k, test[-16], decision.values = TRUE, probability = TRUE)

svm_prediction_k <- predict(svm_classifier_k, test[-16], type = "prob")

confusionMatrix(svm_prediction_k, test$y, mode = "prec_recall")
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  no  yes
##          no 1917 305
##          yes  43 184
##
##           Accuracy : 0.8579
##           95% CI : (0.8434, 0.8715)
##       No Information Rate : 0.8003
##       P-Value [Acc > NIR] : 7.141e-14
##
##           Kappa : 0.4435
##
##  McNemar's Test P-Value : < 2.2e-16
##
##           Precision : 0.8627
##           Recall : 0.9781
##           F1 : 0.9168
##       Prevalence : 0.8003
##       Detection Rate : 0.7828
##       Detection Prevalence : 0.9073
##       Balanced Accuracy : 0.6772
##
##       'Positive' Class : no
##
```


```
# Accuracy : 0.8579
# Precision : 0.8627
# Recall : 0.9781
# F1 : 0.9168
```

```
# AUC-----
prob_svm_k <- attr(pred_svm_k, "probabilities")

colAUC(prob_svm_k, test$y)
```


```
##           no      yes
## no vs. yes 0.8107508 0.8107508
```

```
colAUC(prob_svm_k, test$y, plotROC = "TRUE")
```



```
##           no      yes
## no vs. yes 0.8107508 0.8107508
```

```
auc_svm_k <- roc(response=test$y, predictor=prob_svm_k[,1])
plot(auc_svm_k)
```



```
auc_svm_k$auc # 81.08%
```

```
## Area under the curve: 0.8108
```

```
# Comment: I chose SVM because it is mostly understood when used for binary classification. It is a distance based algorithm, it creates a flat boundary known as hyperplane which divides the space to create fairly homogeneous partitions on either side. It combines kNN and linear regressions. It is very powerful, and model highly complex relationships. SVM model did improve after setting the kernel to radial, as linear did not perform as well. Radial works better here because the underlying data is not linearly separable. After model improvement, the accuracy is 85.79%, F-score is 91.68% and AUC is 81.08% which is the second highest compare to models previously built.
```

2.3. Construction of Decision Trees

```
# Decision Trees:
# Training using data frame without dummy code since decision tree takes both categorical and numeric variables.
decision_tree <- C5.0(training[-11], training$Exited)
summary(decision_tree)
```

```
##
## Call:
## C5.0.default(x = training[-11], y = training$Exited)
##
##
## C5.0 [Release 2.07 GPL Edition]      Wed Apr 15 12:21:41 2020
## -----
##
## Class specified by attribute `outcome'
##
## Read 7350 cases (11 attributes) from undefined.data
##
## Decision tree:
##
## NumOfProducts > 2:
##   ...Balance > 50194.59: yes (113/4)
## :   Balance <= 50194.59:
## :   ...Age > 42: yes (31/2)
## :       Age <= 42:
## :       ...EstimatedSalary <= 153064.9: no (38/13)
## :       EstimatedSalary > 153064.9: yes (10)
## NumOfProducts <= 2:
##   ...Age <= 41: no (4969/480)
##   Age > 41:
##   ...NumOfProducts > 1:
##   ...IsActiveMember = 1: no (532/61)
##   :   IsActiveMember = 0:
##   :   ...Age <= 50: no (289/43)
##   :       Age > 50:
##   :       ...Gender = Female: yes (30/5)
##   :       Gender = Male:
##   :       ...Age <= 53: no (12/3)
##   :       Age > 53: yes (21/4)
##   NumOfProducts <= 1:
##   ...IsActiveMember = 0:
##   ...Age > 47: yes (319/53)
##   :   Age <= 47:
##   :   ...Geography = Germany: yes (109/34)
##   :       Geography in {France,Spain}:
##   :       ...Balance > 97086.4: no (117/36)
##   :       Balance <= 97086.4:
##   :       ...Geography = Spain: yes (48/16)
##   :       Geography = France:
##   :       ...Tenure <= 5: yes (30/9)
##   :       Tenure > 5: no (36/14)
##   IsActiveMember = 1:
##   ...Geography in {France,Spain}: no (467/138)
##   Geography = Germany:
##   ...Balance <= 87347.7: no (15/1)
##   Balance > 87347.7:
##   ...CreditScore > 718: yes (35/6)
##   CreditScore <= 718:
##   ...Gender = Male: no (64/27)
##   Gender = Female:
##   ...HasCrCard = 0: yes (24/7)
##   HasCrCard = 1:
##   ...Balance <= 119565.9: yes (19/4)
##   Balance > 119565.9: no (22/8)
##
##
## Evaluation on training data (7350 cases):
##
##      Decision Tree
##      -----
##      Size      Errors
##
##      23  968(13.2%)  <<
##
##
##      (a)  (b)  <-classified as
##      ----  ----
##      5737  144  (a): class no
##      824   645  (b): class yes
##
##
## Attribute usage:
##
## 100.00% NumOfProducts
```

```
## 98.46% Age
## 29.78% IsActiveMember
## 13.41% Geography
## 8.19% Balance
## 2.61% Gender
## 2.23% CreditScore
## 0.90% Tenure
## 0.88% HasCrCard
## 0.65% EstimatedSalary
##
##
## Time: 0.0 secs
```

```
decision_tree_pred <- predict(decision_tree, testing[-11])

confusionMatrix(testing$Exited, decision_tree_pred, mode = "prec_recall")
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  no  yes
##          no 1898   62
##          yes  278  211
##
##           Accuracy : 0.8612
##           95% CI : (0.8468, 0.8746)
##       No Information Rate : 0.8885
##       P-Value [Acc > NIR] : 1
##
##           Kappa : 0.4793
##
##  Mcnemar's Test P-Value : <2e-16
##
##           Precision : 0.9684
##           Recall : 0.8722
##           F1 : 0.9178
##       Prevalence : 0.8885
##       Detection Rate : 0.7750
##       Detection Prevalence : 0.8003
##       Balanced Accuracy : 0.8226
##
##       'Positive' Class : no
##
```

```
# Accuracy : 0.8612
# Precision : 0.9684
# Recall : 0.8722
# F1: 0.9178

# Tuning Improvement-----
# Trials = 10 is boosting technique, that the algorithm will stop adding trees if the desired overall error rate
# is reached or performance no longer improves with additional of trials

decision_tree_10 <- C5.0(training[-11], training$Exited, trials = 10)

decision_tree_pred_10 <- predict(decision_tree_10, testing[-11])

confusionMatrix(testing$Exited, decision_tree_pred_10, mode = "prec_recall")
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  no  yes
##          no 1869  91
##          yes  262 227
##
##           Accuracy : 0.8559
##           95% CI : (0.8413, 0.8695)
##       No Information Rate : 0.8702
##       P-Value [Acc > NIR] : 0.9825
##
##           Kappa : 0.4809
##
##  McNemar's Test P-Value : <2e-16
##
##           Precision : 0.9536
##           Recall : 0.8771
##           F1 : 0.9137
##       Prevalence : 0.8702
##       Detection Rate : 0.7632
##       Detection Prevalence : 0.8003
##       Balanced Accuracy : 0.7954
##
##       'Positive' Class : no
##
```


```
# Accuracy : 0.8559
# Precision : 0.9536
# Recall : 0.8771
# F1 : 0.9178
# After increasing the number of trials for decision tree, model performance remained the same.
```

```
# AUC-----
pred_dt <- predict(decision_tree_10, testing[-11], type = "prob")

colAUC(pred_dt, test$y)
```


```
##           no      yes
## no vs. yes 0.8401371 0.8401371
```

```
colAUC(pred_dt, test$y, plotROC = TRUE)
```



```
##           no      yes
## no vs. yes 0.8401371 0.8401371
```

```
auc_dt <- roc(response=test$y, predictor=pred_dt[,1])
plot(auc_dt)
```



```
auc_dt$auc # 84.01%
```

```
## Area under the curve: 0.8401
```

Comment: I chose decision tree here because it can handle numeric or categorical features. Decision Trees does n't need dummy codes for categorical variables as it makes if-then branches; like a tree with many branches. And for numerical features, the split is done with the elements higher than a threshold. The accuracy for improved decision tree is 86.12%, F-score is 92%, AUC is 84.01%, this is so far the best performing model.

2.4. Construction of Navie Bayes:

```
set.seed(123)

train_nb <- nb_df[split,]
test_nb <- nb_df[-split,]

m1 <- naiveBayes(train_nb[-11], train_nb$Exited)
pred1 <- predict(m1, test_nb[-11])

confusionMatrix(pred1, test_nb$Exit, mode = "prec_recall")
```



```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  no  yes
##          no 1880 361
##          yes   80 128
##
##           Accuracy : 0.8199
##           95% CI : (0.8041, 0.835)
##       No Information Rate : 0.8003
##       P-Value [Acc > NIR] : 0.007623
##
##           Kappa : 0.2817
##
##  McNemar's Test P-Value : < 2.2e-16
##
##           Precision : 0.8389
##           Recall : 0.9592
##           F1 : 0.8950
##       Prevalence : 0.8003
##       Detection Rate : 0.7677
##       Detection Prevalence : 0.9151
##       Balanced Accuracy : 0.6105
##
##       'Positive' Class : no
##
```

```
# Accuracy : 0.8199
# Precision : 0.8389
# Recall : 0.9592

# Improve model with laplace = 1 -----
m2 <- naiveBayes(train_nb[-11], train_nb$Exited, laplace = 1)
pred2 <- predict(m2, test_nb[-11])

confusionMatrix(pred2, test_nb$Exit, mode = "prec_recall") # 81.99% accuracy rate and kappa score of 0.2817
```


```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  no  yes
##          no 1880 361
##          yes   80 128
##
##           Accuracy : 0.8199
##           95% CI : (0.8041, 0.835)
##       No Information Rate : 0.8003
##       P-Value [Acc > NIR] : 0.007623
##
##           Kappa : 0.2817
##
##  McNemar's Test P-Value : < 2.2e-16
##
##           Precision : 0.8389
##           Recall : 0.9592
##           F1 : 0.8950
##       Prevalence : 0.8003
##       Detection Rate : 0.7677
##       Detection Prevalence : 0.9151
##       Balanced Accuracy : 0.6105
##
##       'Positive' Class : no
##
```

```
# Accuracy : 0.8199
# Precision : 0.8389
# Recall : 0.9592
# F1 : 0.8950

# AUC-----
pred_nb <- predict(m2, test_nb[-11], type = "raw")
colAUC(pred_nb, test$y)
```


```
##                no      yes
## no vs. yes 0.7822863 0.7822863
```

```
colAUC(pred_nb, test$y, plotROC = TRUE)
```



```
##                no      yes
## no vs. yes 0.7822863 0.7822863
```

```
auc_nb <- roc(response=test$y, predictor=pred_nb[,1])
plot(auc_nb)
```



```
auc_nb$auc # 78.23%
```

```
## Area under the curve: 0.7823
```

```
# Comment: I chose Naive Bayes because it is a probabilistic learner that is used for classification. Naive Bayes
accepts only categorical features as it calculates probabilities, so I binned my numerical features. After tuning
the Laplace estimator, the model performance remained the same. Accuracy for Naive Bayes improved model is 81.9
9%, F-score is 89.50% and AUC is 78.23%.
```

CRISP-DM Evaluation

1. Model Evaluation:

- evaluation of fit of models with holdout method
- evaluation with k-fold cross-validation
- tuning of models
- comparison of models
- interpretation of results/prediction with interval
- construction of stacked ensemble model

```
# According to the above holdout methods of 4 methods:
# 1. Decision Trees: Accuracy : 0.8559 Precision : 0.9536 Recall : 0.8771 F1 : 0.9178

# 2. SVM: Accuracy : 0.8579 Precision : 0.8627 Recall : 0.9781 F1 : 0.9168

# 3. Neural Network: Accuracy : 0.7991 Precision : 0.8251 Recall : 0.9505 F1 : 0.8834

# 4. Naive Bayes: Accuracy : 0.8199 Precision : 0.8389 Recall : 0.9592 F1 : 0.8950

# Decision Tree model has the best performance, following by SVM, Neural Network and Naive Bayes.
```

In This section, k-fold cross validation method is used to construct four different models, and each model is tuned to see whether or not it improves performance. Models built are then used to predict test data set and accuracy is calculated.

1. Decision Tree

```
set.seed(1)

# Decision Trees CV and Tuning

# 1. Evaluation of k-folds cross-validation-----
train <- df[split,]
test <- df[-split,]


# Create a control object that uses 10-fold cross validation
ctrl <- trainControl(method="cv", number=10, classProbs = TRUE)

dtFit_cv <- train(Exited ~ ., data = train, method = "C5.0", trControl = ctrl, preProcess = c("center","scale"))

dtFit_cv # Best model is trials = 20, model = rules and winnow = FALSE.
```

```
## C5.0
##
## 7350 samples
## 10 predictor
## 2 classes: 'no', 'yes'
##
## Pre-processing: centered (11), scaled (11)
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 6616, 6615, 6615, 6614, 6615, 6615, ...
## Resampling results across tuning parameters:
##
##  model  winnow  trials  Accuracy  Kappa
##  rules  FALSE   1      0.8583695  0.4833332
##  rules  FALSE  10      0.8590485  0.4859355
##  rules  FALSE  20      0.8606797  0.4860022
##  rules  TRUE   1      0.8564633  0.4730960
##  rules  TRUE  10      0.8589126  0.4750426
##  rules  TRUE  20      0.8583703  0.4770603
##  tree   FALSE   1      0.8580976  0.4857486
##  tree   FALSE  10      0.8608172  0.4848202
##  tree   FALSE  20      0.8598650  0.4885103
##  tree   TRUE   1      0.8563272  0.4729849
##  tree   TRUE  10      0.8556484  0.4733458
##  tree   TRUE  20      0.8578253  0.4827831
##
## Accuracy was used to select the optimal model using the largest value.
## The final values used for the model were trials = 10, model = tree and winnow
## = FALSE.
```

```
plot(dtFit_cv)
```



```
# Predict testing set
p_dt <- predict(dtFit_cv, test[-11])
confusionMatrix(p_dt, test$Exited, mode = "prec_recall")
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  no  yes
##           no 1874 257
##           yes  86 232
##
##           Accuracy : 0.8599
##           95% CI : (0.8456, 0.8735)
##           No Information Rate : 0.8003
##           P-Value [Acc > NIR] : 8.61e-15
##
##           Kappa : 0.4956
##
## Mcnemar's Test P-Value : < 2.2e-16
##
##           Precision : 0.8794
##           Recall : 0.9561
##           F1 : 0.9162
##           Prevalence : 0.8003
##           Detection Rate : 0.7652
##           Detection Prevalence : 0.8702
##           Balanced Accuracy : 0.7153
##
##           'Positive' Class : no
##
```


```
# Accuracy : 0.8599
# Precision : 0.8794
# Recall : 0.9561

# 2. Tuning of model-----
dtGrid <- expand.grid(model="tree", trials = c(1:20), winnow = FALSE)

dtFit_tune <- train(Exited ~ ., data = train,
                    method = "C5.0", metric = "ROC",
                    preProc = c("center", "scale"),
                    trControl = ctrl, tuneGrid = dtGrid)
dtFit_tune # After tuning, best model is trials = 19, model = tree and winnow = FALSE
```

```
## C5.0
##
## 7350 samples
## 10 predictor
## 2 classes: 'no', 'yes'
##
## Pre-processing: centered (11), scaled (11)
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 6615, 6615, 6616, 6615, 6615, 6615, ...
## Resampling results across tuning parameters:
##
## trials Accuracy Kappa
## 1 0.8560575 0.4792764
## 2 0.8548317 0.4623808
## 3 0.8532013 0.4759507
## 4 0.8553780 0.4676654
## 5 0.8557861 0.4924734
## 6 0.8591860 0.4885957
## 7 0.8561943 0.4932071
## 8 0.8600025 0.4924969
## 9 0.8594593 0.5012111
## 10 0.8594574 0.4890303
## 11 0.8593215 0.4997558
## 12 0.8616345 0.4980926
## 13 0.8586418 0.4965608
## 14 0.8612265 0.5005665
## 15 0.8594589 0.5004112
## 16 0.8612280 0.4972572
## 17 0.8591875 0.4964910
## 18 0.8631324 0.5030884
## 19 0.8608207 0.4977209
## 20 0.8619088 0.4988650
##
## Tuning parameter 'model' was held constant at a value of tree
## Tuning
## parameter 'winnow' was held constant at a value of FALSE
## Accuracy was used to select the optimal model using the largest value.
## The final values used for the model were trials = 18, model = tree and winnow
## = FALSE.
```

```
plot(dtFit_tune)
```



```
p_dt_tune <- predict(dtFit_tune, test[-11]) #
confusionMatrix(p_dt_tune, test$Exited, mode = "prec_recall")
```


```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  no  yes
##          no 1875 261
##          yes  85 228
##
##           Accuracy : 0.8587
##           95% CI : (0.8443, 0.8723)
##       No Information Rate : 0.8003
##       P-Value [Acc > NIR] : 3.095e-14
##
##           Kappa : 0.4889
##
##  McNemar's Test P-Value : < 2.2e-16
##
##           Precision : 0.8778
##           Recall : 0.9566
##              F1 : 0.9155
##           Prevalence : 0.8003
##       Detection Rate : 0.7656
##   Detection Prevalence : 0.8722
##       Balanced Accuracy : 0.7114
##
##       'Positive' Class : no
##
```

```
# Accuracy : 0.8587
# Precision : 0.8778
# Recall : 0.9566
# F1 : 0.9162
```

```
# 3. AUC-----
pred_dt_cv <- predict(dtFit_tune, test[-11], type = "prob")
colAUC(pred_dt_cv, test$Exited)
```


```
##           no      yes
## no vs. yes 0.8556592 0.8556592
```

```
colAUC(pred_dt_cv, test$Exited, plotROC = TRUE)
```



```
##           no      yes
## no vs. yes 0.8556592 0.8556592
```

```
auc_dt <- roc(response=test$Exited, predictor=pred_dt_cv[,1])
plot(auc_dt)
```



```
auc_dt$auc # AUC 85.57%
```

```
## Area under the curve: 0.8557
```

```
# Comment: k-fold cross validation did improve the model performance compare to holdout method. Tuning the parameters did not improve the k-fold cross validation method.
```

2. SVM

```
set.seed(1)

# SVM CV and Tuning

# 1. Evaluation of k-folds cross-validation-----
train <- svm_final_df[split,]
test <- svm_final_df[-split,]

# Create a control object that uses 10-fold cross validation
ctrl <- trainControl(method="cv", number=10, classProbs = TRUE)

svmFit_cv <- train(y~ ., data = train, method = "svmRadial",
                  trControl = ctrl, preProcess = c("center","scale"))
```


```
## line search fails -2.224957 -0.4069394 1.132849e-05 -6.518471e-06 -2.945737e-08 -1.491488e-08 -2.364854e-13line
e search fails -2.552532 -0.7434487 1.246619e-05 -8.933555e-06 -3.202841e-08 -1.414955e-08 -2.728664e-13line sear
ch fails -2.286336 -0.4590789 1.100164e-05 -7.368471e-06 -2.595644e-08 -1.090137e-08 -2.052369e-13line search fai
ls -2.604757 -0.7718162 1.498019e-05 -9.761285e-06 -4.581502e-08 -2.433552e-08 -4.487719e-13line search fails -2.
346276 -0.50947 1.274194e-05 -9.079457e-06 -2.915443e-08 -1.124573e-08 -2.693788e-13line search fails -2.561674 -
0.7535846 2.377577e-05 -1.726806e-05 -6.063815e-08 -2.651396e-08 -9.838738e-13line search fails -2.603044 -0.7876
605 1.680958e-05 -1.210135e-05 -4.503906e-08 -2.088394e-08 -5.043637e-13
```

```
svmFit_cv # Cost parameter = 1.00 accuracy = 0.8527891 kappa = 0.4137305
```



```
## Support Vector Machines with Radial Basis Function Kernel
##
## 7350 samples
## 15 predictor
## 2 classes: 'no', 'yes'
##
## Pre-processing: centered (15), scaled (15)
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 6616, 6615, 6615, 6614, 6615, 6615, ...
## Resampling results across tuning parameters:
##
## C      Accuracy   Kappa
## 0.25  0.8553728  0.4644343
## 0.50  0.8507942  0.4344641
## 1.00  0.8548395  0.4432295
##
## Tuning parameter 'sigma' was held constant at a value of 0.04215211
## Accuracy was used to select the optimal model using the largest value.
## The final values used for the model were sigma = 0.04215211 and C = 0.25.
```

```
plot(svmFit_cv)
```



```
# Predict testing set
p_svm <- predict(svmFit_cv, test[-16])
confusionMatrix(p_svm, test$y, mode = "prec_recall")
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  no  yes
##           no 1880 268
##           yes  80 221
##
##           Accuracy : 0.8579
##           95% CI : (0.8434, 0.8715)
##           No Information Rate : 0.8003
##           P-Value [Acc > NIR] : 7.141e-14
##
##           Kappa : 0.4804
##
## Mcnemar's Test P-Value : < 2.2e-16
##
##           Precision : 0.8752
##           Recall : 0.9592
##           F1 : 0.9153
##           Prevalence : 0.8003
##           Detection Rate : 0.7677
##           Detection Prevalence : 0.8771
##           Balanced Accuracy : 0.7056
##
##           'Positive' Class : no
##
```

```
# Accuracy : 0.8579
# Precision : 0.8608
# Recall : 0.9811
```

```
# 2. Tuning of model-----
```

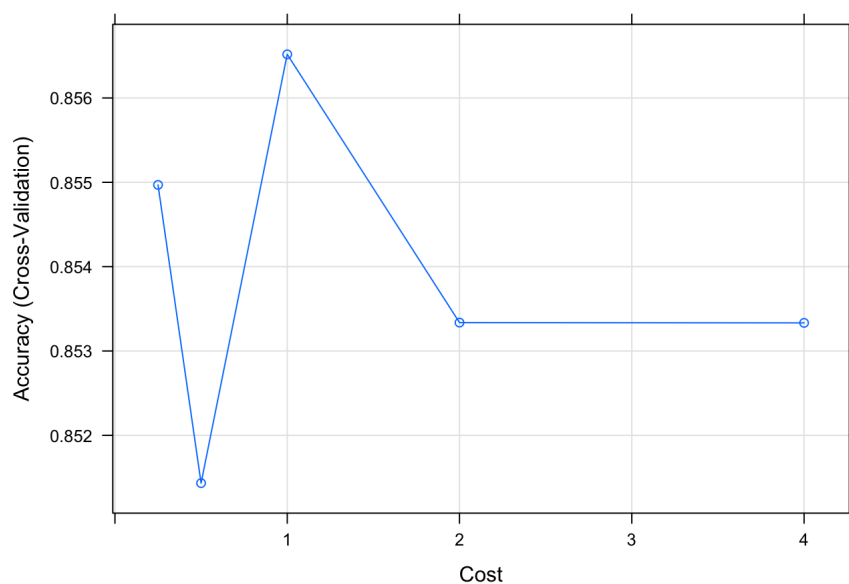
```
svmFit_tune <- train(y ~ ., data = train, method = "svmRadial",
  trControl = ctrl, preProcess = c("center","scale"),
  metric = "ROC", tuneLength = 5, mode = "prec_recall")
```

```
## line search fails -2.557268 -0.7373992 1.588295e-05 -1.121988e-05 -4.262727e-08 -1.987492e-08 -4.540527e-13line
e search fails -2.282037 -0.4563801 1.022401e-05 -6.753859e-06 -2.471233e-08 -1.085317e-08 -1.793582e-13line sear
ch fails -1.942558 -0.1452485 1.966157e-05 -1.289466e-05 -3.492063e-08 -9.710205e-09 -5.613847e-13line search fai
ls -2.522419 -0.7121024 1.806333e-05 -1.261992e-05 -4.753435e-08 -2.196979e-08 -5.813717e-13line search fails -2.
624628 -0.8041854 1.520834e-05 -1.125373e-05 -4.042016e-08 -1.833607e-08 -4.083744e-13line search fails -2.584894
-0.760945 1.436787e-05 -1.020536e-05 -3.907872e-08 -1.835519e-08 -3.741564e-13line search fails -1.918491 -0.1131
688 1.182858e-05 -7.94808e-06 -2.062597e-08 -5.348212e-09 -2.014679e-13line search fails -1.916621 -0.1253765 1.1
74893e-05 -7.818997e-06 -2.048405e-08 -5.434034e-09 -1.981769e-13line search fails -1.876527 -0.10071 1.366534e-0
5 -8.718469e-06 -2.391744e-08 -7.023157e-09 -2.656087e-13line search fails -2.559769 -0.7350673 2.473825e-05 -1.7
84633e-05 -6.369844e-08 -2.767537e-08 -1.081884e-12line search fails -1.877302 -0.08906397 1.200454e-05 -7.721373
e-06 -2.07815e-08 -5.81883e-09 -2.045428e-13
```

```
svmFit_tune # After tuning best model is Cost parameter = 2 accuracy = 0.8552371 kappa = 0.4533225
```

```
## Support Vector Machines with Radial Basis Function Kernel
##
## 7350 samples
## 15 predictor
## 2 classes: 'no', 'yes'
##
## Pre-processing: centered (15), scaled (15)
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 6615, 6614, 6615, 6615, 6615, 6616, ...
## Resampling results across tuning parameters:
##
## C Accuracy Kappa
## 0.25 0.8549693 0.4630186
## 0.50 0.8514333 0.4278080
## 1.00 0.8565174 0.4503513
## 2.00 0.8533366 0.4381417
## 4.00 0.8533333 0.4403914
##
## Tuning parameter 'sigma' was held constant at a value of 0.04320167
## Accuracy was used to select the optimal model using the largest value.
## The final values used for the model were sigma = 0.04320167 and C = 1.
```

```
plot(svmFit_tune)
```



```
# Predict testing set
p_svm_tune <- predict(svmFit_tune, test[-16])
confusionMatrix(p_svm_tune, test$y, mode = "prec_recall")
```


```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  no  yes
##      no  1904  295
##      yes   56  194
##
##           Accuracy : 0.8567
##           95% CI : (0.8422, 0.8703)
##      No Information Rate : 0.8003
##      P-Value [Acc > NIR] : 2.442e-13
##
##           Kappa : 0.4508
##
##  McNemar's Test P-Value : < 2.2e-16
##
##           Precision : 0.8658
##           Recall : 0.9714
##           F1 : 0.9156
##           Prevalence : 0.8003
##           Detection Rate : 0.7775
##           Detection Prevalence : 0.8979
##           Balanced Accuracy : 0.6841
##
##           'Positive' Class : no
##
```

```
# Accuracy : 0.8591
# Precision : 0.8682
# Recall : 0.9714
# F1 : 0.9169
```

```
# 3. AUC-----
pred_svm_cv <- predict(svmFit_tune, test[-16], type = "prob")
colAUC(pred_svm_cv, test$y)
```


```
##           no      yes
## no vs. yes 0.8076583 0.8076583
```

```
colAUC(pred_dt_cv, test$y, plotROC = TRUE)
```



```
##           no      yes
## no vs. yes 0.8556592 0.8556592
```

```
auc_svm <- roc(response=test$y, predictor=pred_svm_cv[,1])
plot(auc_svm)
```



```
auc_svm$auc # AUC 81.07%
```

```
## Area under the curve: 0.8077
```

```
# Comment: k-fold cross validation did not improve the model performance compare to holdout method. Tuning the parameters did improve the k-fold cross validation method.
```

3. Neural Network

```
set.seed(1)

# Neural Network CV and Tuning

# 1. Evaluation of k-folds cross-validation-----
train <- nn_df[split,]
test  <- nn_df[-split,]

# Create a control object that uses 10-fold cross validation
ctrl <- trainControl(method="cv", number=10, classProbs = TRUE)

nnFit_cv <- train(y ~ ., data = train, method = "nnet", metric = "ROC",
                  trControl = ctrl, preProcess = c("center", "scale"))
```

```
## # weights: 9
## initial value 4129.335656
## iter 10 value 2930.582291
## iter 20 value 2838.375747
## iter 30 value 2837.459079
## final value 2837.456903
## converged
## # weights: 25
## initial value 7138.466066
## iter 10 value 2887.821117
## iter 20 value 2699.477166
## iter 30 value 2536.736406
## iter 40 value 2487.233982
## iter 50 value 2483.208641
## iter 60 value 2476.919006
## iter 70 value 2472.068799
## iter 80 value 2469.458414
## iter 90 value 2467.949952
## iter 100 value 2467.435854
## final value 2467.435854
## stopped after 100 iterations
## # weights: 41
## initial value 5415.768682
## iter 10 value 2686.560090
## iter 20 value 2491.268362
## iter 30 value 2454.104489
## iter 40 value 2442.316841
## iter 50 value 2435.365340
## iter 60 value 2431.826271
## iter 70 value 2428.544292
## iter 80 value 2416.979696
## iter 90 value 2411.193580
## iter 100 value 2410.255082
## final value 2410.255082
## stopped after 100 iterations
## # weights: 9
## initial value 3652.866018
## iter 10 value 2948.726328
## iter 20 value 2840.167270
## iter 30 value 2839.195186
## final value 2839.188799
## converged
## # weights: 25
## initial value 4022.642634
## iter 10 value 2843.221443
## iter 20 value 2742.393599
## iter 30 value 2691.087289
## iter 40 value 2631.404544
## iter 50 value 2559.636581
## iter 60 value 2531.271213
## iter 70 value 2516.903350
## iter 80 value 2512.536424
## iter 90 value 2512.132758
## iter 90 value 2512.132738
## iter 90 value 2512.132738
## final value 2512.132738
## converged
## # weights: 41
## initial value 6966.016678
## iter 10 value 2774.146834
## iter 20 value 2685.658828
## iter 30 value 2652.319533
## iter 40 value 2564.712164
## iter 50 value 2494.025988
## iter 60 value 2480.961517
## iter 70 value 2469.749185
## iter 80 value 2456.115604
## iter 90 value 2448.531580
## iter 100 value 2443.341131
## final value 2443.341131
## stopped after 100 iterations
## # weights: 9
## initial value 3807.573169
## iter 10 value 2911.759220
## iter 20 value 2860.022228
## iter 30 value 2843.863823
## iter 40 value 2837.500237
## final value 2837.458793
```

```
## converged
## # weights: 25
## initial value 4043.708589
## iter 10 value 2809.110529
## iter 20 value 2649.235019
## iter 30 value 2525.761946
## iter 40 value 2460.204220
## iter 50 value 2446.303017
## iter 60 value 2445.448676
## iter 70 value 2445.039411
## iter 80 value 2443.701859
## iter 90 value 2438.517086
## iter 100 value 2436.833328
## final value 2436.833328
## stopped after 100 iterations
## # weights: 41
## initial value 3686.784914
## iter 10 value 2833.555137
## iter 20 value 2728.056760
## iter 30 value 2643.434203
## iter 40 value 2573.726484
## iter 50 value 2539.638436
## iter 60 value 2483.899728
## iter 70 value 2459.781317
## iter 80 value 2435.180481
## iter 90 value 2431.223707
## iter 100 value 2426.155556
## final value 2426.155556
## stopped after 100 iterations
## # weights: 9
## initial value 4136.822190
## iter 10 value 3038.282064
## iter 20 value 2842.798707
## iter 30 value 2837.919630
## final value 2837.710409
## converged
## # weights: 25
## initial value 3250.651005
## iter 10 value 2728.247141
## iter 20 value 2500.018368
## iter 30 value 2483.723331
## iter 40 value 2464.915113
## iter 50 value 2452.677426
## iter 60 value 2445.807046
## iter 70 value 2445.237461
## iter 80 value 2444.773276
## iter 90 value 2444.084787
## final value 2444.047777
## converged
## # weights: 41
## initial value 4642.553146
## iter 10 value 2725.852181
## iter 20 value 2602.989018
## iter 30 value 2531.911775
## iter 40 value 2482.363287
## iter 50 value 2463.172116
## iter 60 value 2453.892328
## iter 70 value 2445.998879
## iter 80 value 2444.329091
## iter 90 value 2441.793357
## iter 100 value 2440.730878
## final value 2440.730878
## stopped after 100 iterations
## # weights: 9
## initial value 3448.576188
## iter 10 value 2881.691362
## iter 20 value 2849.253899
## iter 30 value 2839.769003
## final value 2839.462589
## converged
## # weights: 25
## initial value 4938.806793
## iter 10 value 2782.097440
## iter 20 value 2614.213194
## iter 30 value 2554.925018
## iter 40 value 2494.706584
## iter 50 value 2480.133338
## iter 60 value 2475.767156
## iter 70 value 2471.668313
```

```
## iter 80 value 2466.219826
## iter 90 value 2460.458768
## final value 2460.456998
## converged
## # weights: 41
## initial value 6642.997952
## iter 10 value 2755.127809
## iter 20 value 2683.191651
## iter 30 value 2599.300706
## iter 40 value 2524.008540
## iter 50 value 2480.252389
## iter 60 value 2460.658313
## iter 70 value 2453.948946
## iter 80 value 2449.324505
## iter 90 value 2445.563509
## iter 100 value 2442.044874
## final value 2442.044874
## stopped after 100 iterations
## # weights: 9
## initial value 4733.913160
## iter 10 value 2861.884575
## iter 20 value 2840.007993
## iter 30 value 2837.880326
## final value 2837.712196
## converged
## # weights: 25
## initial value 3901.299224
## iter 10 value 2846.999214
## iter 20 value 2794.094910
## iter 30 value 2719.477954
## iter 40 value 2643.392256
## iter 50 value 2597.435622
## iter 60 value 2585.495959
## iter 70 value 2554.693555
## iter 80 value 2537.739832
## iter 90 value 2532.778809
## iter 100 value 2530.951609
## final value 2530.951609
## stopped after 100 iterations
## # weights: 41
## initial value 4944.714405
## iter 10 value 2723.987590
## iter 20 value 2606.571323
## iter 30 value 2544.992034
## iter 40 value 2475.778580
## iter 50 value 2454.701189
## iter 60 value 2446.871713
## iter 70 value 2441.211612
## iter 80 value 2438.191202
## iter 90 value 2437.444823
## iter 100 value 2437.062203
## final value 2437.062203
## stopped after 100 iterations
## # weights: 9
## initial value 4483.646994
## iter 10 value 3059.426147
## iter 20 value 2856.342833
## iter 30 value 2855.685702
## final value 2855.674733
## converged
## # weights: 25
## initial value 5698.713525
## iter 10 value 2846.109864
## iter 20 value 2784.651820
## iter 30 value 2723.286823
## iter 40 value 2655.304930
## iter 50 value 2555.671803
## iter 60 value 2480.600970
## iter 70 value 2464.529679
## iter 80 value 2451.764381
## iter 90 value 2448.817082
## iter 100 value 2448.617113
## final value 2448.617113
## stopped after 100 iterations
## # weights: 41
## initial value 4029.507811
## iter 10 value 2736.988087
## iter 20 value 2592.018651
## iter 30 value 2480.509680
```



```
## iter 40 value 2455.493564
## iter 50 value 2444.872245
## iter 60 value 2443.102473
## iter 70 value 2442.240892
## iter 80 value 2440.319036
## iter 90 value 2437.180595
## iter 100 value 2435.343573
## final value 2435.343573
## stopped after 100 iterations
## # weights: 9
## initial value 5023.298597
## iter 10 value 3203.117053
## iter 20 value 3120.047129
## iter 30 value 2895.309768
## iter 40 value 2860.154270
## iter 50 value 2858.022867
## final value 2858.022290
## converged
## # weights: 25
## initial value 5927.490582
## iter 10 value 2846.657589
## iter 20 value 2792.684268
## iter 30 value 2711.505863
## iter 40 value 2641.426984
## iter 50 value 2595.783759
## iter 60 value 2519.580595
## iter 70 value 2494.403354
## iter 80 value 2469.271074
## iter 90 value 2462.378788
## iter 100 value 2458.323397
## final value 2458.323397
## stopped after 100 iterations
## # weights: 41
## initial value 4998.440499
## iter 10 value 2766.363116
## iter 20 value 2663.964530
## iter 30 value 2614.898030
## iter 40 value 2535.407606
## iter 50 value 2480.888933
## iter 60 value 2469.812194
## iter 70 value 2463.369627
## iter 80 value 2449.456844
## iter 90 value 2444.087551
## iter 100 value 2439.779850
## final value 2439.779850
## stopped after 100 iterations
## # weights: 9
## initial value 4671.548908
## iter 10 value 3299.662962
## iter 20 value 2892.147271
## iter 30 value 2858.207242
## iter 40 value 2855.709171
## final value 2855.676442
## converged
## # weights: 25
## initial value 3369.963783
## iter 10 value 2758.224384
## iter 20 value 2682.273792
## iter 30 value 2617.309088
## iter 40 value 2496.164033
## iter 50 value 2460.251440
## iter 60 value 2455.285669
## iter 70 value 2451.761769
## iter 80 value 2446.757411
## iter 90 value 2445.486664
## iter 100 value 2444.555600
## final value 2444.555600
## stopped after 100 iterations
## # weights: 41
## initial value 5163.008616
## iter 10 value 2685.672610
## iter 20 value 2527.907992
## iter 30 value 2500.691125
## iter 40 value 2483.911708
## iter 50 value 2472.858070
## iter 60 value 2468.932630
## iter 70 value 2466.321105
## iter 80 value 2465.312731
## iter 90 value 2465.142810
```

```
## iter 100 value 2464.478007
## final value 2464.478007
## stopped after 100 iterations
## # weights: 9
## initial value 3542.520783
## iter 10 value 2912.808135
## iter 20 value 2845.475057
## iter 30 value 2842.232920
## final value 2841.955639
## converged
## # weights: 25
## initial value 5576.674481
## iter 10 value 2837.409927
## iter 20 value 2730.501237
## iter 30 value 2715.517631
## iter 40 value 2704.274785
## iter 50 value 2703.011529
## iter 60 value 2686.534727
## iter 70 value 2672.131310
## iter 80 value 2629.528962
## iter 90 value 2611.731855
## iter 100 value 2608.597900
## final value 2608.597900
## stopped after 100 iterations
## # weights: 41
## initial value 3592.917270
## iter 10 value 2703.701047
## iter 20 value 2504.751330
## iter 30 value 2446.151404
## iter 40 value 2431.935395
## iter 50 value 2426.227099
## iter 60 value 2422.317597
## iter 70 value 2420.796230
## iter 80 value 2418.666554
## iter 90 value 2416.134710
## iter 100 value 2414.828001
## final value 2414.828001
## stopped after 100 iterations
## # weights: 9
## initial value 5670.110733
## iter 10 value 3171.079711
## iter 20 value 3082.824500
## iter 30 value 2877.012998
## iter 40 value 2854.850328
## iter 50 value 2844.579184
## final value 2844.472287
## converged
## # weights: 25
## initial value 4110.694710
## iter 10 value 2828.288409
## iter 20 value 2774.974968
## iter 30 value 2725.738472
## iter 40 value 2614.240524
## iter 50 value 2493.571767
## iter 60 value 2483.075273
## iter 70 value 2474.979785
## iter 80 value 2464.739803
## iter 90 value 2460.759371
## iter 100 value 2460.671514
## final value 2460.671514
## stopped after 100 iterations
## # weights: 41
## initial value 3653.488087
## iter 10 value 2721.034327
## iter 20 value 2611.560932
## iter 30 value 2543.796388
## iter 40 value 2490.327822
## iter 50 value 2469.042665
## iter 60 value 2453.545707
## iter 70 value 2440.649555
## iter 80 value 2436.375024
## iter 90 value 2431.783857
## iter 100 value 2427.664211
## final value 2427.664211
## stopped after 100 iterations
## # weights: 9
## initial value 4653.806553
## iter 10 value 2860.250453
## iter 20 value 2844.007376
```

```
## iter 30 value 2841.985830
## final value 2841.957592
## converged
## # weights: 25
## initial value 4327.927093
## iter 10 value 2855.064675
## iter 20 value 2700.484485
## iter 30 value 2654.518981
## iter 40 value 2602.723690
## iter 50 value 2556.604673
## iter 60 value 2535.834192
## iter 70 value 2508.725388
## iter 80 value 2489.298711
## iter 90 value 2482.543171
## iter 100 value 2482.159453
## final value 2482.159453
## stopped after 100 iterations
## # weights: 41
## initial value 5108.281487
## iter 10 value 2740.182308
## iter 20 value 2586.810780
## iter 30 value 2526.373543
## iter 40 value 2474.976860
## iter 50 value 2438.758202
## iter 60 value 2434.068340
## iter 70 value 2432.527477
## iter 80 value 2431.104427
## iter 90 value 2430.071103
## iter 100 value 2429.868731
## final value 2429.868731
## stopped after 100 iterations
## # weights: 9
## initial value 3930.235324
## iter 10 value 3023.055795
## iter 20 value 2856.569405
## iter 30 value 2846.469416
## iter 40 value 2843.424675
## final value 2843.424564
## converged
## # weights: 25
## initial value 4754.282561
## iter 10 value 2789.866657
## iter 20 value 2568.012572
## iter 30 value 2511.517225
## iter 40 value 2504.019831
## iter 50 value 2503.045939
## iter 60 value 2502.965262
## iter 70 value 2502.894060
## iter 80 value 2502.510685
## iter 90 value 2501.386271
## iter 100 value 2470.401129
## final value 2470.401129
## stopped after 100 iterations
## # weights: 41
## initial value 3379.021157
## iter 10 value 2704.852909
## iter 20 value 2606.629758
## iter 30 value 2522.508918
## iter 40 value 2469.244663
## iter 50 value 2441.381156
## iter 60 value 2426.632817
## iter 70 value 2421.694320
## iter 80 value 2418.353655
## iter 90 value 2416.409387
## iter 100 value 2414.382669
## final value 2414.382669
## stopped after 100 iterations
## # weights: 9
## initial value 4070.985456
## iter 10 value 2898.595605
## iter 20 value 2853.643028
## iter 30 value 2846.149035
## final value 2845.083439
## converged
## # weights: 25
## initial value 5968.339753
## iter 10 value 2726.963796
## iter 20 value 2606.657946
## iter 30 value 2542.762373
```

```
## iter 40 value 2503.959403
## iter 50 value 2486.545555
## iter 60 value 2477.668039
## iter 70 value 2471.397401
## iter 80 value 2461.019923
## final value 2460.863964
## converged
## # weights: 41
## initial value 5845.779839
## iter 10 value 2791.966138
## iter 20 value 2693.761835
## iter 30 value 2641.503558
## iter 40 value 2567.930600
## iter 50 value 2507.838400
## iter 60 value 2482.434473
## iter 70 value 2472.123966
## iter 80 value 2466.227329
## iter 90 value 2463.100410
## iter 100 value 2457.589639
## final value 2457.589639
## stopped after 100 iterations
## # weights: 9
## initial value 5157.227551
## iter 10 value 2877.866662
## iter 20 value 2850.043433
## iter 30 value 2843.656312
## final value 2843.426070
## converged
## # weights: 25
## initial value 4343.742921
## iter 10 value 2797.846994
## iter 20 value 2713.319839
## iter 30 value 2648.949035
## iter 40 value 2533.717088
## iter 50 value 2517.089363
## iter 60 value 2508.452358
## iter 70 value 2492.585801
## iter 80 value 2472.044163
## iter 90 value 2460.484448
## iter 100 value 2454.506769
## final value 2454.506769
## stopped after 100 iterations
## # weights: 41
## initial value 4561.579973
## iter 10 value 2780.966264
## iter 20 value 2648.609251
## iter 30 value 2568.384536
## iter 40 value 2482.850245
## iter 50 value 2461.604481
## iter 60 value 2442.372332
## iter 70 value 2434.372900
## iter 80 value 2431.406421
## iter 90 value 2430.409249
## iter 100 value 2429.886538
## final value 2429.886538
## stopped after 100 iterations
## # weights: 9
## initial value 6188.571200
## iter 10 value 3308.806156
## iter 20 value 2943.996730
## iter 30 value 2880.454121
## iter 40 value 2879.802628
## iter 50 value 2877.304074
## iter 60 value 2871.680322
## iter 70 value 2851.175312
## iter 80 value 2837.246167
## final value 2837.128571
## converged
## # weights: 25
## initial value 5335.128292
## iter 10 value 2969.018749
## iter 20 value 2817.469231
## iter 30 value 2800.041950
## iter 40 value 2717.496760
## iter 50 value 2698.530833
## iter 60 value 2696.821318
## iter 70 value 2695.145679
## iter 80 value 2689.035338
## iter 90 value 2686.926952
```

```
## iter 100 value 2686.734580
## final value 2686.734580
## stopped after 100 iterations
## # weights: 41
## initial value 3777.534779
## iter 10 value 2797.690134
## iter 20 value 2704.525125
## iter 30 value 2626.273621
## iter 40 value 2587.672884
## iter 50 value 2580.219650
## iter 60 value 2578.222134
## iter 70 value 2576.068531
## iter 80 value 2555.522175
## iter 90 value 2474.059356
## iter 100 value 2448.686320
## final value 2448.686320
## stopped after 100 iterations
## # weights: 9
## initial value 6232.098899
## iter 10 value 3292.063378
## iter 20 value 3144.917102
## iter 30 value 2857.389206
## iter 40 value 2847.149385
## iter 50 value 2839.436079
## final value 2839.423400
## converged
## # weights: 25
## initial value 6140.751103
## iter 10 value 2853.440913
## iter 20 value 2665.566852
## iter 30 value 2552.814676
## iter 40 value 2496.340866
## iter 50 value 2477.380503
## iter 60 value 2474.882124
## iter 70 value 2470.249685
## iter 80 value 2456.678574
## iter 90 value 2442.445614
## iter 100 value 2441.901852
## final value 2441.901852
## stopped after 100 iterations
## # weights: 41
## initial value 4890.470845
## iter 10 value 2824.683406
## iter 20 value 2598.361095
## iter 30 value 2513.785913
## iter 40 value 2481.307807
## iter 50 value 2464.298631
## iter 60 value 2456.213222
## iter 70 value 2448.511878
## iter 80 value 2444.873165
## iter 90 value 2443.265533
## iter 100 value 2441.654500
## final value 2441.654500
## stopped after 100 iterations
## # weights: 9
## initial value 5695.796395
## iter 10 value 2891.287866
## iter 20 value 2846.613179
## iter 30 value 2838.667751
## final value 2837.130181
## converged
## # weights: 25
## initial value 3423.032490
## iter 10 value 2821.954110
## iter 20 value 2766.983034
## iter 30 value 2733.482677
## iter 40 value 2708.154198
## iter 50 value 2691.942012
## iter 60 value 2691.412243
## iter 70 value 2690.980239
## iter 80 value 2689.086355
## iter 90 value 2688.756975
## iter 100 value 2688.612101
## final value 2688.612101
## stopped after 100 iterations
## # weights: 41
## initial value 6051.599586
## iter 10 value 2730.375667
## iter 20 value 2600.620385
```

```
## iter 30 value 2525.258441
## iter 40 value 2466.129231
## iter 50 value 2441.216664
## iter 60 value 2431.203973
## iter 70 value 2421.797344
## iter 80 value 2414.999365
## iter 90 value 2412.494363
## iter 100 value 2404.850655
## final value 2404.850655
## stopped after 100 iterations
## # weights: 9
## initial value 4250.148786
## iter 10 value 2867.393100
## iter 20 value 2848.510512
## iter 30 value 2847.956120
## final value 2847.940768
## converged
## # weights: 25
## initial value 3790.239713
## iter 10 value 2829.767990
## iter 20 value 2722.205668
## iter 30 value 2692.386466
## iter 40 value 2636.485367
## iter 50 value 2594.167048
## iter 60 value 2588.940396
## iter 70 value 2581.237882
## iter 80 value 2541.973296
## iter 90 value 2505.262128
## iter 100 value 2475.860931
## final value 2475.860931
## stopped after 100 iterations
## # weights: 41
## initial value 6049.954476
## iter 10 value 2768.114458
## iter 20 value 2664.793723
## iter 30 value 2599.664938
## iter 40 value 2540.459300
## iter 50 value 2512.694080
## iter 60 value 2482.359481
## iter 70 value 2468.517963
## iter 80 value 2462.717495
## iter 90 value 2461.455399
## iter 100 value 2460.825276
## final value 2460.825276
## stopped after 100 iterations
## # weights: 9
## initial value 4846.561175
## iter 10 value 2875.009088
## iter 20 value 2850.736250
## iter 30 value 2849.593979
## final value 2849.504600
## converged
## # weights: 25
## initial value 4587.337501
## iter 10 value 2811.701746
## iter 20 value 2761.035848
## iter 30 value 2569.332353
## iter 40 value 2505.354702
## iter 50 value 2475.406226
## iter 60 value 2465.751139
## iter 70 value 2460.058753
## iter 80 value 2453.830435
## iter 90 value 2448.631917
## final value 2448.597901
## converged
## # weights: 41
## initial value 4051.600476
## iter 10 value 2901.052816
## iter 20 value 2840.884154
## iter 30 value 2787.151488
## iter 40 value 2748.830929
## iter 50 value 2667.799295
## iter 60 value 2518.068963
## iter 70 value 2461.729057
## iter 80 value 2444.781666
## iter 90 value 2438.483531
## iter 100 value 2437.693287
## final value 2437.693287
## stopped after 100 iterations
```

```
## # weights: 9
## initial value 3981.259217
## iter 10 value 2999.011200
## iter 20 value 2883.990040
## iter 30 value 2855.377388
## iter 40 value 2847.944697
## final value 2847.943082
## converged
## # weights: 25
## initial value 4403.397886
## iter 10 value 2697.748609
## iter 20 value 2600.860503
## iter 30 value 2520.932098
## iter 40 value 2490.737583
## iter 50 value 2484.386213
## iter 60 value 2481.499073
## iter 70 value 2474.517364
## iter 80 value 2464.451473
## iter 90 value 2461.410685
## iter 100 value 2457.408957
## final value 2457.408957
## stopped after 100 iterations
## # weights: 41
## initial value 5101.576453
## iter 10 value 2825.115509
## iter 20 value 2662.945763
## iter 30 value 2623.036929
## iter 40 value 2600.207760
## iter 50 value 2575.755528
## iter 60 value 2526.692743
## iter 70 value 2459.444579
## iter 80 value 2432.056293
## iter 90 value 2427.482210
## iter 100 value 2424.613345
## final value 2424.613345
## stopped after 100 iterations
## # weights: 9
## initial value 5077.536745
## iter 10 value 2876.440821
## iter 20 value 2835.115193
## iter 30 value 2833.787973
## final value 2833.777650
## converged
## # weights: 25
## initial value 4530.051715
## iter 10 value 2846.633831
## iter 20 value 2824.508377
## iter 30 value 2798.115459
## iter 40 value 2719.183589
## iter 50 value 2619.373944
## iter 60 value 2593.983067
## iter 70 value 2571.149733
## iter 80 value 2528.630984
## iter 90 value 2482.996095
## iter 100 value 2479.185031
## final value 2479.185031
## stopped after 100 iterations
## # weights: 41
## initial value 3386.212384
## iter 10 value 2720.231545
## iter 20 value 2610.777351
## iter 30 value 2556.416452
## iter 40 value 2503.524821
## iter 50 value 2468.596839
## iter 60 value 2455.934705
## iter 70 value 2449.733732
## iter 80 value 2445.803987
## iter 90 value 2443.955131
## iter 100 value 2442.176572
## final value 2442.176572
## stopped after 100 iterations
## # weights: 9
## initial value 5476.547420
## iter 10 value 2949.957680
## iter 20 value 2860.101808
## iter 30 value 2839.236844
## iter 40 value 2835.401671
## final value 2835.401095
## converged
```

```
## # weights: 25
## initial value 4172.466851
## iter 10 value 2767.096003
## iter 20 value 2691.995046
## iter 30 value 2602.185748
## iter 40 value 2522.303586
## iter 50 value 2484.541672
## iter 60 value 2476.639052
## iter 70 value 2471.835470
## iter 80 value 2467.922010
## iter 90 value 2465.330142
## final value 2465.311899
## converged
## # weights: 41
## initial value 3364.005642
## iter 10 value 2813.350943
## iter 20 value 2729.238907
## iter 30 value 2689.219810
## iter 40 value 2633.600427
## iter 50 value 2524.561391
## iter 60 value 2492.456339
## iter 70 value 2479.806462
## iter 80 value 2475.088391
## iter 90 value 2466.327178
## iter 100 value 2460.224145
## final value 2460.224145
## stopped after 100 iterations
## # weights: 9
## initial value 3703.415383
## iter 10 value 2861.509391
## iter 20 value 2834.544055
## iter 30 value 2833.842595
## final value 2833.779289
## converged
## # weights: 25
## initial value 4784.664880
## iter 10 value 2899.887097
## iter 20 value 2731.027008
## iter 30 value 2699.817755
## iter 40 value 2590.389655
## iter 50 value 2544.979488
## iter 60 value 2504.236977
## iter 70 value 2494.099677
## iter 80 value 2482.651700
## iter 90 value 2480.938444
## iter 100 value 2480.273051
## final value 2480.273051
## stopped after 100 iterations
## # weights: 41
## initial value 4783.193647
## iter 10 value 2714.633926
## iter 20 value 2542.871939
## iter 30 value 2504.998801
## iter 40 value 2471.582690
## iter 50 value 2450.880740
## iter 60 value 2445.953293
## iter 70 value 2441.400565
## iter 80 value 2439.362804
## iter 90 value 2438.629533
## iter 100 value 2438.280524
## final value 2438.280524
## stopped after 100 iterations
## # weights: 9
## initial value 3681.662888
## iter 10 value 2900.830729
## iter 20 value 2873.866051
## iter 30 value 2859.735196
## iter 40 value 2843.791260
## iter 50 value 2843.677087
## final value 2843.676829
## converged
## # weights: 25
## initial value 4593.404519
## iter 10 value 2680.757234
## iter 20 value 2571.614005
## iter 30 value 2477.579927
## iter 40 value 2447.535968
## iter 50 value 2442.203302
## iter 60 value 2440.010719
```



```
## iter 70 value 2438.686024
## iter 80 value 2436.640165
## iter 90 value 2435.403289
## iter 100 value 2435.148569
## final value 2435.148569
## stopped after 100 iterations
## # weights: 41
## initial value 4055.726032
## iter 10 value 2837.577158
## iter 20 value 2789.847271
## iter 30 value 2697.873015
## iter 40 value 2645.109060
## iter 50 value 2605.660410
## iter 60 value 2520.282615
## iter 70 value 2444.694365
## iter 80 value 2425.867521
## iter 90 value 2423.273204
## iter 100 value 2421.728418
## final value 2421.728418
## stopped after 100 iterations
## # weights: 9
## initial value 3937.009537
## iter 10 value 2940.849652
## iter 20 value 2845.995728
## iter 30 value 2845.358416
## final value 2845.344174
## converged
## # weights: 25
## initial value 6832.980453
## iter 10 value 2888.305476
## iter 20 value 2842.648245
## iter 30 value 2830.534149
## iter 40 value 2784.208570
## iter 50 value 2615.187936
## iter 60 value 2512.280910
## iter 70 value 2489.733996
## iter 80 value 2476.971001
## iter 90 value 2473.772780
## iter 100 value 2472.680577
## final value 2472.680577
## stopped after 100 iterations
## # weights: 41
## initial value 3867.544390
## iter 10 value 2780.692691
## iter 20 value 2694.018093
## iter 30 value 2664.285927
## iter 40 value 2630.174913
## iter 50 value 2540.164281
## iter 60 value 2475.210835
## iter 70 value 2458.596980
## iter 80 value 2448.694179
## iter 90 value 2445.555703
## iter 100 value 2442.018055
## final value 2442.018055
## stopped after 100 iterations
## # weights: 9
## initial value 5558.164908
## iter 10 value 2877.875185
## iter 20 value 2845.374275
## iter 30 value 2843.693695
## final value 2843.678518
## converged
## # weights: 25
## initial value 7530.850050
## iter 10 value 2822.407941
## iter 20 value 2689.734751
## iter 30 value 2612.262544
## iter 40 value 2544.154132
## iter 50 value 2503.701765
## iter 60 value 2498.660571
## iter 70 value 2485.076115
## iter 80 value 2473.141139
## iter 90 value 2471.401041
## iter 100 value 2469.358396
## final value 2469.358396
## stopped after 100 iterations
## # weights: 41
## initial value 3789.129651
## iter 10 value 2778.227077
```

```
## iter 20 value 2644.128597
## iter 30 value 2587.170320
## iter 40 value 2570.346601
## iter 50 value 2563.498742
## iter 60 value 2549.593288
## iter 70 value 2538.285002
## iter 80 value 2531.587801
## iter 90 value 2528.743160
## iter 100 value 2526.475452
## final value 2526.475452
## stopped after 100 iterations
## # weights: 9
## initial value 4502.261827
## iter 10 value 2908.987173
## iter 20 value 2860.313442
## final value 2854.274894
## converged
## # weights: 25
## initial value 4945.139341
## iter 10 value 2749.175677
## iter 20 value 2671.836821
## iter 30 value 2562.119746
## iter 40 value 2483.278625
## iter 50 value 2467.366322
## iter 60 value 2464.128881
## iter 70 value 2458.719612
## iter 80 value 2452.958654
## iter 90 value 2451.916864
## iter 100 value 2451.769798
## final value 2451.769798
## stopped after 100 iterations
## # weights: 41
## initial value 4293.305775
## iter 10 value 2704.648890
## iter 20 value 2582.135307
## iter 30 value 2498.092797
## iter 40 value 2450.311318
## iter 50 value 2431.726871
## iter 60 value 2424.895535
## iter 70 value 2419.857940
## iter 80 value 2417.483726
## iter 90 value 2415.235534
## iter 100 value 2414.504232
## final value 2414.504232
## stopped after 100 iterations
## # weights: 9
## initial value 4211.727693
## iter 10 value 2876.236783
## iter 20 value 2856.110623
## iter 30 value 2855.911321
## final value 2855.910512
## converged
## # weights: 25
## initial value 4264.108321
## iter 10 value 2838.895383
## iter 20 value 2786.334180
## iter 30 value 2654.110624
## iter 40 value 2619.207529
## iter 50 value 2574.081040
## iter 60 value 2545.470477
## iter 70 value 2517.141787
## iter 80 value 2488.349298
## iter 90 value 2472.714333
## iter 100 value 2465.694522
## final value 2465.694522
## stopped after 100 iterations
## # weights: 41
## initial value 4457.737004
## iter 10 value 2780.611174
## iter 20 value 2656.822358
## iter 30 value 2601.930981
## iter 40 value 2512.730038
## iter 50 value 2481.258386
## iter 60 value 2464.633476
## iter 70 value 2459.862163
## iter 80 value 2456.408478
## iter 90 value 2453.539502
## iter 100 value 2446.443204
## final value 2446.443204
```

```
## stopped after 100 iterations
## # weights: 9
## initial value 5407.968335
## iter 10 value 2915.722376
## iter 20 value 2859.210854
## iter 30 value 2854.834048
## final value 2854.276551
## converged
## # weights: 25
## initial value 5144.602684
## iter 10 value 2837.164902
## iter 20 value 2750.754857
## iter 30 value 2719.353997
## iter 40 value 2697.058960
## iter 50 value 2689.927673
## iter 60 value 2688.873843
## iter 70 value 2688.513365
## iter 80 value 2686.519015
## iter 90 value 2671.732025
## iter 100 value 2669.533095
## final value 2669.533095
## stopped after 100 iterations
## # weights: 41
## initial value 6362.315721
## iter 10 value 2777.985787
## iter 20 value 2668.125472
## iter 30 value 2628.196005
## iter 40 value 2579.808157
## iter 50 value 2531.874453
## iter 60 value 2468.540939
## iter 70 value 2428.692417
## iter 80 value 2421.941620
## iter 90 value 2421.705072
## iter 100 value 2421.465059
## final value 2421.465059
## stopped after 100 iterations
## # weights: 41
## initial value 4227.361483
## iter 10 value 3121.219838
## iter 20 value 2981.837188
## iter 30 value 2898.723770
## iter 40 value 2773.941063
## iter 50 value 2741.172714
## iter 60 value 2730.796474
## iter 70 value 2723.283799
## iter 80 value 2721.920580
## iter 90 value 2721.333267
## iter 100 value 2720.718885
## final value 2720.718885
## stopped after 100 iterations
```

nnFit_cv # Best model is size = 5 and decay = 1e-04 accuracy 84.08%

```
## Neural Network
##
## 7350 samples
## 6 predictor
## 2 classes: 'no', 'yes'
##
## Pre-processing: centered (6), scaled (6)
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 6616, 6615, 6615, 6614, 6615, ...
## Resampling results across tuning parameters:
##
## size decay Accuracy Kappa
## 1 0e+00 0.7917002 0.1631688
## 1 1e-04 0.7917002 0.1631688
## 1 1e-01 0.7912924 0.1598192
## 3 0e+00 0.8342881 0.3986396
## 3 1e-04 0.8326543 0.3900103
## 3 1e-01 0.8397310 0.4195130
## 5 0e+00 0.8408192 0.4293501
## 5 1e-04 0.8416344 0.4219078
## 5 1e-01 0.8390496 0.4127871
##
## Accuracy was used to select the optimal model using the largest value.
## The final values used for the model were size = 5 and decay = 1e-04.
```

```
# Predict testing set
p_nn <- predict(nnFit_cv, test[-7])
confusionMatrix(p_nn, test$y, mode = "prec_recall")
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  no  yes
##          no 1822 279
##          yes  138 210
##
##           Accuracy : 0.8297
##           95% CI : (0.8142, 0.8444)
##       No Information Rate : 0.8003
##       P-Value [Acc > NIR] : 0.0001166
##
##           Kappa : 0.4026
##
##  Mcnemar's Test P-Value : 7.09e-12
##
##           Precision : 0.8672
##           Recall : 0.9296
##           F1 : 0.8973
##           Prevalence : 0.8003
##       Detection Rate : 0.7440
##       Detection Prevalence : 0.8579
##       Balanced Accuracy : 0.6795
##
##       'Positive' Class : no
##
```

```
# Accuracy : 0.8297
# Precision : 0.8672
# Recall : 0.9296
```

```
# 2. Tuning of model-----
nnetGrid <- expand.grid(size = seq(from = 1, to = 3, by = 1), # Number of nodes
                      decay = seq(from = 0.1, to = 0.2, by = 0.1))

nnFit_tune <- train(y ~ ., data = train, method = "nnet",
                  trControl = ctrl, preProcess = c("center", "scale"),
                  metric = "ROC", tuneGrid = nnetGrid)
```

```
## # weights: 9
## initial value 4669.851282
## iter 10 value 2942.089462
## iter 20 value 2846.709175
## iter 30 value 2844.190166
## final value 2844.134343
## converged
## # weights: 17
## initial value 5675.706082
## iter 10 value 2818.893897
## iter 20 value 2747.591529
## iter 30 value 2695.233697
## iter 40 value 2674.264840
## iter 50 value 2650.310505
## iter 60 value 2591.621596
## iter 70 value 2567.863068
## iter 80 value 2565.121605
## iter 90 value 2563.503152
## final value 2563.411744
## converged
## # weights: 25
## initial value 3679.396826
## iter 10 value 2875.366913
## iter 20 value 2833.899351
## iter 30 value 2814.600348
## iter 40 value 2806.192777
## iter 50 value 2716.754473
## iter 60 value 2513.563792
## iter 70 value 2472.982839
## iter 80 value 2453.602616
## iter 90 value 2450.103720
## iter 100 value 2447.559552
## final value 2447.559552
## stopped after 100 iterations
## # weights: 9
## initial value 4030.930246
## iter 10 value 2860.282545
## iter 20 value 2845.750554
## iter 30 value 2845.742443
## final value 2845.740352
## converged
## # weights: 17
## initial value 5325.744925
## iter 10 value 2911.683102
## iter 20 value 2794.563715
## iter 30 value 2759.406879
## iter 40 value 2755.793646
## iter 50 value 2755.528728
## final value 2755.509552
## converged
## # weights: 25
## initial value 4803.277691
## iter 10 value 2846.251396
## iter 20 value 2822.503053
## iter 30 value 2770.599753
## iter 40 value 2747.887941
## iter 50 value 2743.955555
## iter 60 value 2740.154479
## iter 70 value 2737.952239
## iter 80 value 2736.707485
## iter 90 value 2733.523162
## iter 100 value 2731.331571
## final value 2731.331571
## stopped after 100 iterations
## # weights: 9
## initial value 4421.109344
## iter 10 value 2902.415848
## iter 20 value 2851.779898
## iter 30 value 2842.935142
## final value 2841.366446
## converged
## # weights: 17
## initial value 5590.934959
## iter 10 value 2795.894165
## iter 20 value 2749.126552
## iter 30 value 2713.424079
## iter 40 value 2677.834096
## iter 50 value 2658.646319
```

```
## iter 60 value 2627.712424
## iter 70 value 2574.352668
## iter 80 value 2569.391752
## iter 90 value 2565.254621
## final value 2565.054511
## converged
## # weights: 25
## initial value 5789.807376
## iter 10 value 2748.766662
## iter 20 value 2636.596646
## iter 30 value 2531.370099
## iter 40 value 2491.183834
## iter 50 value 2467.536906
## iter 60 value 2462.547125
## iter 70 value 2454.010578
## iter 80 value 2445.845394
## iter 90 value 2445.182259
## final value 2445.179267
## converged
## # weights: 9
## initial value 4242.557936
## iter 10 value 2870.917881
## iter 20 value 2842.991701
## final value 2842.964333
## converged
## # weights: 17
## initial value 5947.145020
## iter 10 value 2879.587766
## iter 20 value 2838.299027
## iter 30 value 2825.343246
## iter 40 value 2761.521618
## iter 50 value 2711.505254
## iter 60 value 2667.604071
## iter 70 value 2649.882857
## iter 80 value 2627.602102
## iter 90 value 2597.211448
## iter 100 value 2595.292261
## final value 2595.292261
## stopped after 100 iterations
## # weights: 25
## initial value 7762.914695
## iter 10 value 2936.756529
## iter 20 value 2852.935640
## iter 30 value 2813.657422
## iter 40 value 2659.825302
## iter 50 value 2603.915004
## iter 60 value 2577.779413
## iter 70 value 2535.938017
## iter 80 value 2502.766804
## iter 90 value 2491.849095
## iter 100 value 2489.876045
## final value 2489.876045
## stopped after 100 iterations
## # weights: 9
## initial value 3506.992659
## iter 10 value 3162.108076
## iter 20 value 2879.778931
## iter 30 value 2856.768103
## iter 40 value 2850.472913
## final value 2850.442901
## converged
## # weights: 17
## initial value 3906.451252
## iter 10 value 2929.136256
## iter 20 value 2854.037452
## iter 30 value 2826.521966
## iter 40 value 2747.541780
## iter 50 value 2741.769309
## iter 60 value 2677.422890
## iter 70 value 2623.196148
## iter 80 value 2613.347380
## iter 90 value 2596.233157
## iter 100 value 2595.164025
## final value 2595.164025
## stopped after 100 iterations
## # weights: 25
## initial value 3649.029284
## iter 10 value 2885.111670
## iter 20 value 2801.240388
```

```
## iter 30 value 2744.840735
## iter 40 value 2715.937160
## iter 50 value 2701.134730
## iter 60 value 2700.038697
## iter 70 value 2603.299372
## iter 80 value 2548.556656
## iter 90 value 2525.869512
## iter 100 value 2495.770427
## final value 2495.770427
## stopped after 100 iterations
## # weights: 9
## initial value 5701.901058
## iter 10 value 2896.311345
## iter 20 value 2857.953741
## iter 30 value 2852.183503
## final value 2852.036888
## converged
## # weights: 17
## initial value 3489.533721
## iter 10 value 2784.908983
## iter 20 value 2748.537157
## iter 30 value 2693.316790
## iter 40 value 2672.916862
## iter 50 value 2654.997019
## iter 60 value 2622.922271
## iter 70 value 2620.441640
## final value 2620.434974
## converged
## # weights: 25
## initial value 4297.966890
## iter 10 value 2881.127793
## iter 20 value 2796.737844
## iter 30 value 2714.538034
## iter 40 value 2592.128629
## iter 50 value 2533.620423
## iter 60 value 2517.461428
## iter 70 value 2502.465620
## iter 80 value 2485.086060
## iter 90 value 2475.537376
## iter 100 value 2473.383465
## final value 2473.383465
## stopped after 100 iterations
## # weights: 9
## initial value 3481.569908
## iter 10 value 2869.388382
## iter 20 value 2832.003958
## iter 30 value 2831.770194
## final value 2831.769764
## converged
## # weights: 17
## initial value 4640.418509
## iter 10 value 2961.455365
## iter 20 value 2838.285902
## iter 30 value 2816.138362
## iter 40 value 2705.661947
## iter 50 value 2661.791454
## iter 60 value 2612.586063
## iter 70 value 2561.009631
## iter 80 value 2557.535940
## iter 90 value 2551.486976
## final value 2551.417908
## converged
## # weights: 25
## initial value 4120.418285
## iter 10 value 2731.493719
## iter 20 value 2647.886529
## iter 30 value 2608.176971
## iter 40 value 2571.630431
## iter 50 value 2518.352312
## iter 60 value 2472.183282
## iter 70 value 2445.665455
## iter 80 value 2435.974928
## iter 90 value 2433.232353
## final value 2433.224107
## converged
## # weights: 9
## initial value 5499.047685
## iter 10 value 2936.086368
## iter 20 value 2844.541465
```

```
## iter 30 value 2833.433308
## final value 2833.418686
## converged
## # weights: 17
## initial value 3498.749569
## iter 10 value 2853.423840
## iter 20 value 2834.092942
## iter 30 value 2811.504555
## iter 40 value 2750.031202
## iter 50 value 2565.662009
## iter 60 value 2534.299486
## iter 70 value 2529.481564
## iter 80 value 2525.742816
## final value 2525.616829
## converged
## # weights: 25
## initial value 4672.602071
## iter 10 value 2769.048158
## iter 20 value 2569.486497
## iter 30 value 2480.506982
## iter 40 value 2456.932361
## iter 50 value 2451.119154
## iter 60 value 2447.956324
## iter 70 value 2446.535971
## iter 80 value 2446.033413
## final value 2446.016640
## converged
## # weights: 9
## initial value 4674.370551
## iter 10 value 2950.966285
## iter 20 value 2875.572508
## iter 30 value 2871.950786
## final value 2871.629993
## converged
## # weights: 17
## initial value 6834.284516
## iter 10 value 2908.547863
## iter 20 value 2833.100926
## iter 30 value 2825.584692
## iter 40 value 2825.312032
## iter 50 value 2813.812619
## iter 60 value 2771.443407
## iter 70 value 2741.444389
## final value 2741.287408
## converged
## # weights: 25
## initial value 3292.068647
## iter 10 value 2738.199155
## iter 20 value 2685.278120
## iter 30 value 2654.987084
## iter 40 value 2613.261744
## iter 50 value 2580.249885
## iter 60 value 2560.141257
## iter 70 value 2540.147832
## iter 80 value 2529.851048
## iter 90 value 2511.670547
## iter 100 value 2499.389943
## final value 2499.389943
## stopped after 100 iterations
## # weights: 9
## initial value 3767.735859
## iter 10 value 2905.454464
## iter 20 value 2849.256647
## iter 30 value 2844.532136
## final value 2844.480104
## converged
## # weights: 17
## initial value 5805.913204
## iter 10 value 2959.436097
## iter 20 value 2852.284110
## iter 30 value 2776.638515
## iter 40 value 2759.547047
## iter 50 value 2759.154195
## final value 2759.099974
## converged
## # weights: 25
## initial value 5269.087507
## iter 10 value 2774.597057
## iter 20 value 2704.110612
```



```
## iter 30 value 2676.722447
## iter 40 value 2615.115834
## iter 50 value 2566.887360
## iter 60 value 2548.470568
## iter 70 value 2527.369956
## iter 80 value 2511.139962
## iter 90 value 2504.631649
## iter 100 value 2501.460888
## final value 2501.460888
## stopped after 100 iterations
## # weights: 9
## initial value 4920.537551
## iter 10 value 2886.659340
## iter 20 value 2855.726071
## iter 30 value 2851.573530
## final value 2851.523075
## converged
## # weights: 17
## initial value 4984.528352
## iter 10 value 2841.805501
## iter 20 value 2718.565114
## iter 30 value 2677.477486
## iter 40 value 2662.880495
## iter 50 value 2660.601955
## iter 60 value 2642.084216
## iter 70 value 2614.344525
## iter 80 value 2602.022869
## iter 90 value 2587.877177
## iter 100 value 2583.461672
## final value 2583.461672
## stopped after 100 iterations
## # weights: 25
## initial value 3555.716881
## iter 10 value 2807.504599
## iter 20 value 2748.387972
## iter 30 value 2694.086947
## iter 40 value 2554.529392
## iter 50 value 2502.833680
## iter 60 value 2493.617313
## iter 70 value 2486.405539
## iter 80 value 2478.289103
## iter 90 value 2468.729010
## iter 100 value 2466.192964
## final value 2466.192964
## stopped after 100 iterations
## # weights: 9
## initial value 6216.096833
## iter 10 value 3021.855973
## iter 20 value 2857.448246
## iter 30 value 2853.193028
## final value 2853.174348
## converged
## # weights: 17
## initial value 3801.746968
## iter 10 value 2869.041621
## iter 20 value 2832.071868
## iter 30 value 2820.207015
## iter 40 value 2607.067358
## iter 50 value 2571.473586
## iter 60 value 2551.320705
## final value 2551.303348
## converged
## # weights: 25
## initial value 5594.089998
## iter 10 value 2821.662637
## iter 20 value 2775.321763
## iter 30 value 2710.877390
## iter 40 value 2590.450921
## iter 50 value 2519.544745
## iter 60 value 2501.867233
## iter 70 value 2488.262258
## iter 80 value 2479.511563
## iter 90 value 2474.155648
## iter 100 value 2471.951516
## final value 2471.951516
## stopped after 100 iterations
## # weights: 9
## initial value 5869.306657
## iter 10 value 3275.237394
```

```
## iter 20 value 3130.224031
## iter 30 value 3094.877441
## iter 40 value 2868.412903
## iter 50 value 2839.856159
## iter 60 value 2836.115160
## iter 70 value 2836.103939
## final value 2836.102748
## converged
## # weights: 17
## initial value 5540.967200
## iter 10 value 2826.203428
## iter 20 value 2686.240005
## iter 30 value 2644.992098
## iter 40 value 2638.875177
## iter 50 value 2629.033571
## iter 60 value 2616.427500
## iter 70 value 2596.949557
## iter 80 value 2588.055893
## iter 90 value 2561.759433
## iter 100 value 2558.957807
## final value 2558.957807
## stopped after 100 iterations
## # weights: 25
## initial value 5185.785078
## iter 10 value 2819.848929
## iter 20 value 2758.957855
## iter 30 value 2701.724157
## iter 40 value 2652.758035
## iter 50 value 2618.733143
## iter 60 value 2603.259710
## iter 70 value 2578.409571
## iter 80 value 2496.548058
## iter 90 value 2481.368905
## iter 100 value 2465.897151
## final value 2465.897151
## stopped after 100 iterations
## # weights: 9
## initial value 4658.210330
## iter 10 value 2853.502441
## iter 20 value 2837.188005
## final value 2837.187056
## converged
## # weights: 17
## initial value 4708.026454
## iter 10 value 2942.558714
## iter 20 value 2775.492872
## iter 30 value 2707.341097
## iter 40 value 2672.882034
## iter 50 value 2644.312764
## iter 60 value 2603.176958
## iter 70 value 2595.733888
## iter 80 value 2595.296504
## iter 90 value 2595.186346
## iter 90 value 2595.186337
## iter 90 value 2595.186337
## final value 2595.186337
## converged
## # weights: 25
## initial value 5907.699412
## iter 10 value 2885.386912
## iter 20 value 2779.288525
## iter 30 value 2638.501049
## iter 40 value 2538.927147
## iter 50 value 2525.454119
## iter 60 value 2519.627827
## iter 70 value 2512.566175
## iter 80 value 2508.854302
## iter 90 value 2507.467579
## final value 2507.463278
## converged
## # weights: 9
## initial value 7380.801823
## iter 10 value 2868.520279
## iter 20 value 2846.715256
## iter 30 value 2843.165097
## final value 2843.129646
## converged
## # weights: 17
## initial value 3841.951686
```

```
## iter 10 value 2817.272051
## iter 20 value 2752.526361
## iter 30 value 2711.922898
## iter 40 value 2663.880869
## iter 50 value 2647.273119
## iter 60 value 2590.800206
## iter 70 value 2579.314711
## iter 80 value 2578.655730
## iter 90 value 2578.196485
## final value 2578.193419
## converged
## # weights: 25
## initial value 4559.385720
## iter 10 value 2850.908453
## iter 20 value 2766.217925
## iter 30 value 2705.865101
## iter 40 value 2678.596122
## iter 50 value 2652.121155
## iter 60 value 2639.537740
## iter 70 value 2622.319385
## iter 80 value 2583.413209
## iter 90 value 2553.871445
## iter 100 value 2546.767623
## final value 2546.767623
## stopped after 100 iterations
## # weights: 9
## initial value 4128.983546
## iter 10 value 2914.105206
## iter 20 value 2852.970737
## iter 30 value 2845.024521
## final value 2844.800035
## converged
## # weights: 17
## initial value 3768.927157
## iter 10 value 2869.946230
## iter 20 value 2840.914035
## iter 30 value 2804.181997
## iter 40 value 2797.115667
## iter 50 value 2752.311837
## final value 2752.307019
## converged
## # weights: 25
## initial value 4080.121318
## iter 10 value 2773.493503
## iter 20 value 2719.502014
## iter 30 value 2648.036361
## iter 40 value 2561.675593
## iter 50 value 2525.302870
## iter 60 value 2507.737690
## iter 70 value 2495.460385
## iter 80 value 2480.828034
## iter 90 value 2466.224826
## iter 100 value 2465.917993
## final value 2465.917993
## stopped after 100 iterations
## # weights: 9
## initial value 3863.920967
## iter 10 value 2896.039415
## iter 20 value 2861.445152
## iter 30 value 2855.949639
## final value 2855.607300
## converged
## # weights: 17
## initial value 6044.052559
## iter 10 value 2878.317839
## iter 20 value 2846.612678
## iter 30 value 2837.561426
## iter 40 value 2768.349918
## iter 50 value 2754.303988
## iter 60 value 2749.486554
## iter 60 value 2749.486529
## iter 60 value 2749.486529
## final value 2749.486529
## converged
## # weights: 25
## initial value 3506.614915
## iter 10 value 2861.347853
## iter 20 value 2797.623010
## iter 30 value 2652.175269
```

```
## iter 40 value 2586.233785
## iter 50 value 2496.718538
## iter 60 value 2473.880368
## iter 70 value 2467.942559
## iter 80 value 2460.016956
## iter 90 value 2455.546215
## iter 100 value 2455.534659
## final value 2455.534659
## stopped after 100 iterations
## # weights: 9
## initial value 3749.610590
## iter 10 value 3105.445082
## iter 20 value 2857.734659
## iter 30 value 2857.279664
## final value 2857.266461
## converged
## # weights: 17
## initial value 4366.652930
## iter 10 value 2858.571725
## iter 20 value 2830.633095
## iter 30 value 2821.758154
## iter 40 value 2817.035049
## iter 50 value 2816.156848
## iter 60 value 2814.824852
## final value 2814.824735
## converged
## # weights: 25
## initial value 3759.266951
## iter 10 value 2816.771080
## iter 20 value 2593.968293
## iter 30 value 2530.811334
## iter 40 value 2507.828368
## iter 50 value 2502.404613
## iter 60 value 2483.250696
## iter 70 value 2475.698256
## iter 80 value 2469.429112
## iter 90 value 2468.740548
## final value 2468.740456
## converged
## # weights: 9
## initial value 5662.081047
## iter 10 value 2964.418701
## iter 20 value 2853.558165
## iter 30 value 2853.001007
## final value 2852.985689
## converged
## # weights: 17
## initial value 4158.276093
## iter 10 value 2860.113060
## iter 20 value 2780.309626
## iter 30 value 2712.627218
## iter 40 value 2681.868194
## iter 50 value 2666.209069
## iter 60 value 2629.384472
## iter 70 value 2594.522634
## iter 80 value 2590.014270
## iter 90 value 2588.440151
## final value 2588.415250
## converged
## # weights: 25
## initial value 4348.577269
## iter 10 value 2840.668874
## iter 20 value 2808.012753
## iter 30 value 2646.600445
## iter 40 value 2558.967016
## iter 50 value 2520.428494
## iter 60 value 2509.574425
## iter 70 value 2501.478097
## iter 80 value 2487.376044
## iter 90 value 2468.713821
## iter 100 value 2467.362368
## final value 2467.362368
## stopped after 100 iterations
## # weights: 9
## initial value 3601.319992
## iter 10 value 2904.911451
## iter 20 value 2856.271452
## iter 30 value 2853.877476
## final value 2853.787695
```

```

## converged
## # weights: 17
## initial value 4954.619328
## iter 10 value 2898.844410
## iter 20 value 2814.572723
## iter 30 value 2730.394644
## iter 40 value 2695.464033
## iter 50 value 2675.204669
## iter 60 value 2637.822069
## iter 70 value 2619.569915
## iter 80 value 2618.492792
## iter 90 value 2618.176657
## final value 2618.176008
## converged
## # weights: 25
## initial value 5540.846638
## iter 10 value 2870.760042
## iter 20 value 2779.966159
## iter 30 value 2685.675736
## iter 40 value 2548.582023
## iter 50 value 2519.672972
## iter 60 value 2509.920595
## iter 70 value 2495.528496
## iter 80 value 2481.916183
## iter 90 value 2476.575566
## iter 100 value 2476.433131
## final value 2476.433131
## stopped after 100 iterations
## # weights: 25
## initial value 5754.220814
## iter 10 value 3112.469003
## iter 20 value 3014.412880
## iter 30 value 2942.419279
## iter 40 value 2843.059573
## iter 50 value 2769.850795
## iter 60 value 2754.535783
## iter 70 value 2745.078361
## iter 80 value 2729.817868
## iter 90 value 2726.053647
## iter 100 value 2725.940830
## final value 2725.940830
## stopped after 100 iterations

```


```
nnFit_tune # After tuning, best model is size = 3 and decay = 0.2 Accuracy is 83.86% and Kappa = 0.4097141
```

```

## Neural Network
##
## 7350 samples
## 6 predictor
## 2 classes: 'no', 'yes'
##
## Pre-processing: centered (6), scaled (6)
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 6615, 6615, 6614, 6615, 6615, ...
## Resampling results across tuning parameters:
##
##  size decay Accuracy Kappa
##  1    0.1  0.7945535  0.1806682
##  1    0.2  0.7940093  0.1706728
##  2    0.1  0.8282959  0.3610426
##  2    0.2  0.8243472  0.3467579
##  3    0.1  0.8410876  0.4176904
##  3    0.2  0.8386370  0.4086754
##
## Accuracy was used to select the optimal model using the largest value.
## The final values used for the model were size = 3 and decay = 0.1.

```

```
plot(svmFit_tune)
```



```
# Predict testing set
p_nn_tune <- predict(nnFit_tune, test[-7])
confusionMatrix(p_nn_tune, test$y, mode = "prec_recall")
```


```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  no  yes
##      no  1841  288
##      yes   119  201
##
##           Accuracy : 0.8338
##           95% CI : (0.8185, 0.8484)
##      No Information Rate : 0.8003
##      P-Value [Acc > NIR] : 1.289e-05
##
##           Kappa : 0.4025
##
##  Mcnemar's Test P-Value : < 2.2e-16
##
##           Precision : 0.8647
##           Recall : 0.9393
##           F1 : 0.9005
##           Prevalence : 0.8003
##           Detection Rate : 0.7517
##           Detection Prevalence : 0.8693
##           Balanced Accuracy : 0.6752
##
##           'Positive' Class : no
##
```

```
# Accuracy : 0.8338
# Precision : 0.8647
# Recall : 0.9393
```

```
# 3. AUC-----
pred_nn_cv <- predict(nnFit_tune, test[-7], type = "prob")
colAUC(pred_nn_cv, test$y)
```


```
##           no      yes
## no vs. yes 0.8162514 0.8162514
```

```
colAUC(pred_nn_cv, test$y, plotROC = TRUE)
```



```
##           no      yes
## no vs. yes 0.8162514 0.8162514
```

```
auc_nn <- roc(response=test$y, predictor=pred_nn_cv[,1])
plot(auc_dt)
```



```
auc_nn$auc # AUC 81.63%
```

```
## Area under the curve: 0.8163
```

```
# Comment: k-fold cross validation did improve the model performance compare to holdout method. Tuning the parameters did improve the holdout method.
```

4. Naive Bayes

```
# Naive Bayes CV and Tuning

# 1. Evaluation of k-folds cross-validation-----

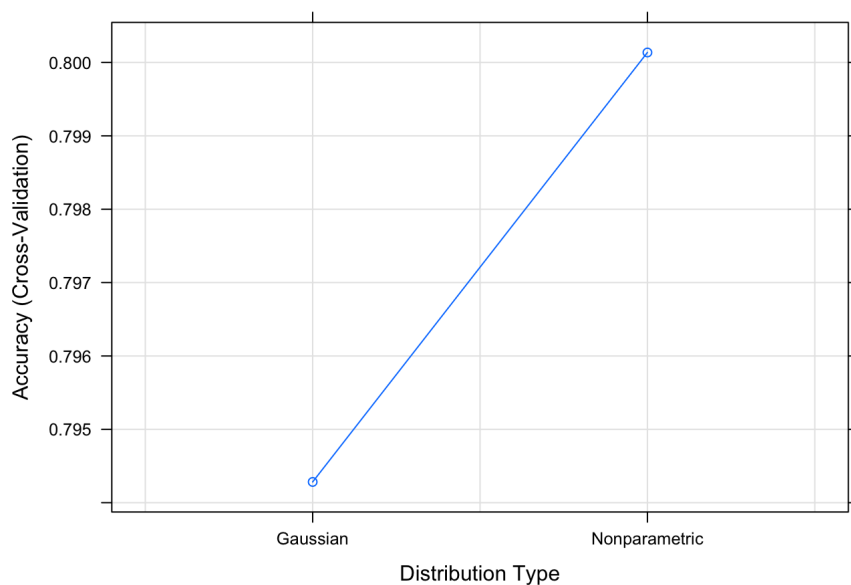
train <- nb_df[split,]
test <- nb_df[-split,]

nbFit_cv <- train(Exited ~ ., data = train, method = "nb", trControl = ctrl)

nbFit_cv
```

```
## Naive Bayes
##
## 7350 samples
## 10 predictor
## 2 classes: 'no', 'yes'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 6616, 6615, 6615, 6615, 6615, ...
## Resampling results across tuning parameters:
##
## usekernel Accuracy Kappa
## FALSE 0.7942832 0.3020774
## TRUE 0.8001362 0.0000000
##
## Tuning parameter 'fL' was held constant at a value of 0
## Tuning
## parameter 'adjust' was held constant at a value of 1
## Accuracy was used to select the optimal model using the largest value.
## The final values used for the model were fL = 0, usekernel = TRUE and adjust
## = 1.
```

```
plot(nbFit_cv)
```



```
# Predict testing set
p_nb <- predict(nbFit_cv, test[-11])
confusionMatrix(p_nb, test$Exited, mode = "prec_recall")
```



```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  no  yes
##           no 1960 489
##           yes    0    0
##
##           Accuracy : 0.8003
##           95% CI : (0.7839, 0.816)
##           No Information Rate : 0.8003
##           P-Value [Acc > NIR] : 0.5121
##
##           Kappa : 0
##
## Mcnemar's Test P-Value : <2e-16
##
##           Precision : 0.8003
##           Recall : 1.0000
##           F1 : 0.8891
##           Prevalence : 0.8003
##           Detection Rate : 0.8003
##           Detection Prevalence : 1.0000
##           Balanced Accuracy : 0.5000
##
##           'Positive' Class : no
##
```

```
# Accuracy : 0.8003
# Precision : 0.8003
# Recall : 1.0000
```

```
# Comment: Naive Bayes did not perform well, it has a sensitivity of 0, I did not perform a tuning and cross validation for it as it did not improve anything and taking a very long time to run. The cross validation result and holdout method produced the same result.
```

Evaluation of k-fold cross validation, not surprisingly, Decision Tree model again out performed the other models with 86% accuracy, precision of 87.78%, recall of 95.66% and F-score of 91.62% and AUC of 85.57%, follow by SVM with accuracy of 85.91%, precision of 86.82%, recall of 97.14%, F-score of 91.69% and AUC of 81.07%, follow by Neural Network with 83.38%, precision of 86.47%, recall of 93.93% and F-score of 90.05% and AUC of 81.63%. Naive Byes again performed the worst. I will exclude it from the Stacking Ensemble model.

In this section, Stacking Ensemble of top three performing models: Decision Tree, SVM and Neural Network are stacked to build an ensemble model. And Bagging Ensemble for Decision

Tree is built. The Goal is to compare bagging ensemble of Decision Tree and Stacked Ensemble model of Decision Trees, SVM and Neural Network

Construction of Stacking Ensemble Model and AUC Evaluation

```
set.seed(1)

train <- df[split,]
test <- df[-split,]

# Example of Stacking algorithms
# create submodels
control <- trainControl(method="cv", number=5, savePredictions = "final", classProbs=TRUE)

algorithmList <- c('C5.0', 'svmRadial', 'nnet')

# Stacking Algorithms - Run multiple algorithms in one call.
models <- caretList(Exited ~., data=train, trControl=control, methodList=algorithmList)
```

```
## line search fails -2.560131 -0.7218277 1.040596e-05 -6.232883e-06 -3.786098e-08 -2.16508e-08 -2.590329e-13# weights: 14
## initial value 3644.996684
## final value 2922.374580
## converged
## # weights: 40
## initial value 3042.681029
## iter 10 value 2897.353445
## iter 20 value 2896.545037
## iter 30 value 2896.524189
## final value 2896.523700
## converged
## # weights: 66
## initial value 3956.641029
## iter 10 value 2897.845234
## iter 20 value 2897.370780
## final value 2897.359498
## converged
## # weights: 14
## initial value 3594.996252
## iter 10 value 2942.556452
## iter 20 value 2903.124056
## iter 30 value 2802.351299
## iter 40 value 2688.789559
## iter 50 value 2585.433677
## iter 60 value 2456.801272
## iter 70 value 2415.422725
## iter 80 value 2408.653611
## iter 90 value 2406.189914
## final value 2406.076553
## converged
## # weights: 40
## initial value 6922.481628
## iter 10 value 2942.170228
## iter 20 value 2938.568118
## iter 30 value 2897.663185
## iter 40 value 2896.310266
## iter 50 value 2894.704713
## iter 60 value 2868.921241
## iter 70 value 2744.913200
## iter 80 value 2707.366383
## iter 90 value 2672.420209
## iter 100 value 2669.988795
## final value 2669.988795
## stopped after 100 iterations
## # weights: 66
## initial value 3059.716257
## iter 10 value 2903.108124
## final value 2900.275296
## converged
## # weights: 14
## initial value 3020.482264
## final value 2942.366658
## converged
## # weights: 40
## initial value 3396.398731
## iter 10 value 2905.901811
## final value 2895.037700
## converged
## # weights: 66
## initial value 3489.015481
## iter 10 value 2938.891534
## final value 2938.888284
## converged
## # weights: 14
## initial value 4471.454353
## final value 2940.979424
## converged
## # weights: 40
## initial value 3485.456067
## iter 10 value 2903.666979
## iter 20 value 2900.176643
## iter 30 value 2899.707962
## iter 40 value 2899.644293
## iter 50 value 2899.595296
## final value 2899.595164
## converged
## # weights: 66
```

```
## initial value 5139.002600
## final value 2940.979424
## converged
## # weights: 14
## initial value 3398.520357
## iter 10 value 2905.683090
## iter 20 value 2903.821013
## final value 2903.814374
## converged
## # weights: 40
## initial value 5933.291285
## iter 10 value 2999.459516
## iter 20 value 2941.222852
## iter 30 value 2933.861241
## iter 40 value 2900.090032
## iter 50 value 2896.165971
## iter 60 value 2895.388141
## iter 70 value 2879.302045
## iter 80 value 2860.200670
## iter 90 value 2762.175198
## iter 100 value 2750.227587
## final value 2750.227587
## stopped after 100 iterations
## # weights: 66
## initial value 9814.353956
## iter 10 value 2902.830951
## iter 20 value 2902.022182
## iter 30 value 2901.973593
## iter 40 value 2897.083478
## iter 50 value 2892.539006
## iter 60 value 2828.026695
## iter 70 value 2771.060146
## iter 80 value 2756.051876
## iter 90 value 2755.001520
## iter 100 value 2753.953727
## final value 2753.953727
## stopped after 100 iterations
## # weights: 14
## initial value 4453.572955
## final value 2940.979798
## converged
## # weights: 40
## initial value 4743.659085
## iter 10 value 2925.884673
## final value 2901.227129
## converged
## # weights: 66
## initial value 3225.629826
## iter 10 value 2901.644784
## iter 20 value 2897.524766
## iter 30 value 2897.133476
## iter 40 value 2896.846511
## iter 50 value 2896.646811
## iter 60 value 2896.325216
## iter 70 value 2896.063990
## iter 80 value 2895.939517
## final value 2895.938691
## converged
## # weights: 14
## initial value 3307.919521
## final value 2940.979424
## converged
## # weights: 40
## initial value 3582.993964
## iter 10 value 2936.981994
## iter 20 value 2936.959984
## final value 2936.959785
## converged
## # weights: 66
## initial value 3150.839490
## iter 10 value 2908.861631
## iter 20 value 2907.020058
## iter 30 value 2906.744522
## iter 40 value 2905.865464
## final value 2905.849733
## converged
## # weights: 14
## initial value 3084.784896
## iter 10 value 2941.088651
```

```
## iter 10 value 2941.088639
## iter 10 value 2941.088634
## final value 2941.088634
## converged
## # weights: 40
## initial value 7342.636765
## iter 10 value 2934.503102
## iter 20 value 2928.761469
## iter 30 value 2927.512639
## iter 40 value 2927.270994
## iter 50 value 2913.815272
## iter 60 value 2905.282376
## iter 70 value 2904.311754
## iter 80 value 2904.257959
## final value 2904.078250
## converged
## # weights: 66
## initial value 2941.094622
## iter 10 value 2912.372729
## iter 20 value 2904.052933
## iter 30 value 2897.930149
## iter 40 value 2881.891947
## iter 50 value 2864.852810
## iter 60 value 2812.586510
## iter 70 value 2764.857582
## iter 80 value 2726.898889
## iter 90 value 2698.289311
## iter 100 value 2688.800469
## final value 2688.800469
## stopped after 100 iterations
## # weights: 14
## initial value 4798.486783
## final value 2940.979769
## converged
## # weights: 40
## initial value 3172.091126
## iter 10 value 2926.005300
## iter 20 value 2923.329500
## iter 30 value 2909.276929
## iter 40 value 2908.929334
## iter 50 value 2907.421637
## iter 60 value 2907.393902
## final value 2907.312832
## converged
## # weights: 66
## initial value 3910.697235
## iter 10 value 2931.379348
## iter 20 value 2929.547403
## iter 30 value 2929.525974
## iter 40 value 2928.879485
## iter 50 value 2928.581339
## iter 60 value 2928.505069
## iter 70 value 2928.484266
## iter 80 value 2928.157634
## final value 2928.157236
## converged
## # weights: 14
## initial value 3816.249265
## final value 2940.979424
## converged
## # weights: 40
## initial value 4229.004201
## iter 10 value 2931.257031
## iter 10 value 2931.257027
## iter 10 value 2931.257027
## final value 2931.257027
## converged
## # weights: 66
## initial value 3456.969424
## iter 10 value 2903.596157
## iter 20 value 2892.667844
## iter 30 value 2892.580766
## iter 40 value 2892.066596
## iter 50 value 2891.813048
## iter 60 value 2891.747960
## iter 70 value 2891.491882
## iter 70 value 2891.491882
## iter 70 value 2891.491882
## final value 2891.491882
```


```
## converged
## # weights: 14
## initial value 4128.272542
## iter 10 value 2940.344782
## iter 20 value 2940.003172
## iter 30 value 2939.732658
## iter 30 value 2939.732633
## iter 40 value 2939.354448
## iter 40 value 2939.354429
## iter 40 value 2939.354427
## final value 2939.354427
## converged
## # weights: 40
## initial value 5242.026392
## iter 10 value 2895.021968
## iter 20 value 2894.204671
## iter 30 value 2892.959206
## iter 40 value 2889.707163
## iter 50 value 2889.544350
## final value 2889.541411
## converged
## # weights: 66
## initial value 4120.969455
## iter 10 value 2930.379945
## iter 20 value 2904.584228
## iter 30 value 2898.169523
## iter 40 value 2892.886544
## iter 50 value 2891.945692
## iter 60 value 2891.860758
## iter 70 value 2890.915752
## iter 80 value 2890.730379
## iter 90 value 2890.318082
## iter 100 value 2889.692294
## final value 2889.692294
## stopped after 100 iterations
## # weights: 14
## initial value 5111.676160
## iter 10 value 2940.108268
## final value 2940.099144
## converged
## # weights: 40
## initial value 8655.979186
## iter 10 value 2940.732543
## iter 20 value 2939.621109
## final value 2939.619237
## converged
## # weights: 66
## initial value 3721.099381
## iter 10 value 2929.494101
## iter 20 value 2896.819140
## iter 30 value 2896.253020
## iter 40 value 2895.973889
## iter 50 value 2895.807613
## final value 2895.678205
## converged
## # weights: 14
## initial value 4461.250215
## final value 2940.979424
## converged
## # weights: 40
## initial value 8045.000867
## iter 10 value 2940.095573
## final value 2940.087378
## converged
## # weights: 66
## initial value 5669.278246
## iter 10 value 2929.807161
## iter 20 value 2929.180878
## iter 30 value 2928.993005
## iter 40 value 2928.895767
## final value 2928.892117
## converged
## # weights: 14
## initial value 2970.777418
## iter 10 value 2932.336053
## iter 20 value 2903.788120
## iter 30 value 2903.655832
## final value 2903.655124
## converged
```

```
## # weights: 40
## initial value 6809.701187
## iter 10 value 2941.502759
## iter 20 value 2890.968844
## iter 30 value 2883.721835
## iter 40 value 2845.222264
## iter 50 value 2827.150530
## iter 60 value 2766.097676
## iter 70 value 2757.340247
## iter 80 value 2744.478330
## iter 90 value 2738.626405
## iter 100 value 2734.778802
## final value 2734.778802
## stopped after 100 iterations
## # weights: 66
## initial value 4646.346690
## iter 10 value 2894.837941
## iter 20 value 2888.452789
## iter 30 value 2886.802723
## iter 40 value 2886.470761
## iter 50 value 2876.624431
## iter 60 value 2871.546277
## iter 70 value 2778.563734
## iter 80 value 2769.393214
## iter 90 value 2763.155713
## iter 100 value 2721.370290
## final value 2721.370290
## stopped after 100 iterations
## # weights: 14
## initial value 3516.356942
## final value 2939.763891
## converged
## # weights: 40
## initial value 3448.522018
## iter 10 value 2888.871140
## iter 10 value 2888.871140
## iter 10 value 2888.871140
## final value 2888.871140
## converged
## # weights: 66
## initial value 3580.646919
## iter 10 value 2926.850259
## iter 20 value 2925.757302
## iter 30 value 2899.848433
## final value 2892.984883
## converged
## # weights: 14
## initial value 4487.084012
## iter 10 value 3676.686843
## final value 3676.685651
## converged
```

```
results <- resamples(models)
summary(results)
```

```
##
## Call:
## summary.resamples(object = results)
##
## Models: C5.0, svmRadial, nnet
## Number of resamples: 5
##
## Accuracy
##           Min.   1st Qu.   Median     Mean   3rd Qu.   Max. NA's
## C5.0       0.8544218 0.8564626 0.8578231 0.8594558 0.8639456 0.8646259    0
## svmRadial 0.8469388 0.8503401 0.8578231 0.8551020 0.8585034 0.8619048    0
## nnet      0.8000000 0.8000000 0.8000000 0.8031293 0.8000000 0.8156463    0
##
## Kappa
##           Min.   1st Qu.   Median     Mean   3rd Qu.   Max. NA's
## C5.0       0.4649258 0.4709651 0.4723866 0.48417873 0.5057129 0.5069034    0
## svmRadial 0.4375593 0.4410569 0.4649258 0.46465353 0.4794795 0.5002462    0
## nnet      0.0000000 0.0000000 0.0000000 0.05397797 0.0000000 0.2698899    0
```

```
dotplot(results)
```




Comment: As we can see decision tree has the highest performance in the ensemble model.

```
ensemble_1 <- caretEnsemble(models,
  metric = "ROC",
  trControl = control)
summary(ensemble_1)
```

```
## The following models were ensembled: C5.0, svmRadial, nnet
## They were weighted:
## 2.7757 -5.4308 -0.671 0.5265
## The resulting Accuracy is: 0.8607
## The fit for each individual model on the Accuracy is:
##   method Accuracy AccuracySD
##   C5.0 0.8594558 0.004578590
##   svmRadial 0.8551020 0.006216213
##   nnet 0.8031293 0.006997220
```

```
plot(ensemble_1)
```




```
# From the plot, we can see that C5.0 has the best performance.

# Combine the predictions of multiple models to form a final prediction.

# Ensemble the predictions of `models` to form a new combined prediction based on glm.
stack.glm <- caretStack(models, method = "glm", metric="Accuracy", trControl=control)

print(stack.glm)
```

```
## A glm ensemble of 3 base models: C5.0, svmRadial, nnet
##
## Ensemble results:
## Generalized Linear Model
##
## 7350 samples
## 3 predictor
## 2 classes: 'no', 'yes'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold)
## Summary of sample sizes: 5880, 5880, 5880, 5880, 5880
## Resampling results:
##
## Accuracy Kappa
## 0.8601361 0.4928501
```

```
pred_ensemble <- predict(stack.glm, test[-11])

confusionMatrix(pred_ensemble, test$Exited, mode = "prec_recall")
```


```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  no  yes
##           no 1869 252
##           yes 91 237
##
##           Accuracy : 0.8599
##           95% CI : (0.8456, 0.8735)
##           No Information Rate : 0.8003
##           P-Value [Acc > NIR] : 8.61e-15
##
##           Kappa : 0.5
##
## Mcnemar's Test P-Value : < 2.2e-16
##
##           Precision : 0.8812
##           Recall : 0.9536
##           F1 : 0.9160
##           Prevalence : 0.8003
##           Detection Rate : 0.7632
##           Detection Prevalence : 0.8661
##           Balanced Accuracy : 0.7191
##
##           'Positive' Class : no
##
```

```
# Accuracy : 0.8599
# Precision : 0.8812
# Recall : 0.9536
# F1 : 0.9160
```

```
# AUC-----
pred_stack <- predict(stack.glm, test[-11], type = "prob")
colAUC(pred_stack, test$Exited)
```


```
##           [,1]
## no vs. yes 0.8570531
```

```
colAUC(pred_stack, test$Exited, plotROC = TRUE)
```



```
##           [,1]
## no vs. yes 0.8570531
```

```
auc_nb <- roc(response=test$Exited, predictor=pred_stack)
plot(auc_nb)
```



```
auc_nb$auc # 85.71%
```

```
## Area under the curve: 0.8571
```

```
# Comment: My stacked ensemble improved performance on average, it is the best performing model.
```

Use Bootstrap and aggregating: Decision Tree Bagging and AUC Evaluation

Bagging: generates a number of training dataset by bootstrap sampling the original training data. These data sets then then used to generate a set of models using a single learning algorithm, here will be Decision Tree. The models' predictions are combined using voting for classification or averaging for numeric

prediction. Similar to bagging, boosting also resamples the training data, except it generates complimentary learners and assigns weights, the ones with better performance have greater weights over the ensemble's final prediction.

```
# Top three strong learners
# Bagging and boosting works well with decision tree models, here I chose bagging.
set.seed(1)
ctrl <- trainControl(method = "cv", number = 5) # by number of decision trees voting in the ensemble
bag <- train(Exited ~ ., data = training, method = "treebag", trControl = ctrl)

pred_bag <- predict(bag, testing[-11])
confusionMatrix(pred_bag, testing$Exited, mode = "prec_recall")
```


```
## Confusion Matrix and Statistics
##
##              Reference
## Prediction    no  yes
##      no  1846  247
##      yes   114  242
##
##              Accuracy : 0.8526
##              95% CI : (0.8379, 0.8664)
##      No Information Rate : 0.8003
##      P-Value [Acc > NIR] : 1.191e-11
##
##              Kappa : 0.4864
##
##  McNemar's Test P-Value : 3.722e-12
##
##              Precision : 0.8820
##              Recall : 0.9418
##              F1 : 0.9109
##              Prevalence : 0.8003
##              Detection Rate : 0.7538
##      Detection Prevalence : 0.8546
##              Balanced Accuracy : 0.7184
##
##              'Positive' Class : no
##
```

```
# Accuracy : 0.8526
# Precision : 0.8820
# Recall : 0.9418
# F1: 0.9109
```

```
# AUC-----
pred_bag_auc <- predict(bag, testing[-11], type = "prob")
colAUC(pred_bag_auc, testing$Exited)
```


```
##              no      yes
## no vs. yes 0.8228194 0.8228194
```

```
colAUC(pred_bag_auc, testing$Exited, plotROC = TRUE)
```



```
##          no      yes
## no vs. yes 0.8228194 0.8228194
```

```
auc_bag <- roc(response=testing$Exited, predictor=pred_bag_auc[,1])
plot(auc_bag)
```



```
auc_bag$auc # AUC 82.28%
```

```
## Area under the curve: 0.8228
```

In conclusion, Stacked Ensemble model performed the best, it has an accuracy of 86% precision of 88.12%, recall of 95.36% and F-score of 91.60% and AUC of 85.71%. Decision tree model alone which is very close to the stacked ensemble model with accuracy of 86.12% and precision of 96.84% and recall of 87.22% and F-score of 91.78% and AUC of 84.01%, is the second best performing model. In my case, depends on the computer and its heuristic running method, both Stacked Ensemble model and Decision Tree base model alone performed equally

well. Out of the four models built, Naive Bayes performed the worst, a lesson I learned here is perhaps kmeans should be used for clustering before binning to get the the optimal bin boundaries.

CRISP-DM Deployment

- According to the model prediction, about 86.12% accuracy, banks could use this prediction to predict whether a customer will churn or not. 86.12% will correctly predict churn movement.
- Bank could save money in the long run by targeting which type of customers to keep.
- Helps bank to develop loyalty programs and retention campaigns to keep as many customers as possible.

References

Machine Learning with R, Third Edition, Brett Lantz

<https://www.dataquest.io/blog/top-10-machine-learning-algorithms-for-beginners/> (<https://www.dataquest.io/blog/top-10-machine-learning-algorithms-for-beginners/>)

<http://www.sthda.com/english/articles/38-regression-model-validation/157-cross-validation-essentials-in-r/> (<http://www.sthda.com/english/articles/38-regression-model-validation/157-cross-validation-essentials-in-r/>)

<https://topepo.github.io/caret/train-models-by-tag.html#boosting> (<https://topepo.github.io/caret/train-models-by-tag.html#boosting>)

https://github.com/suesui117/da5030_final_project/blob/main/DA5030.Proj.Sui.html#boosting

https://www.saedsayad.com/k_nearest_neighbors.htm (https://www.saedsayad.com/k_nearest_neighbors.htm)

<https://medium.com/@eijaz/holdout-vs-cross-validation-in-machine-learning-7637112d3f8f> (<https://medium.com/@eijaz/holdout-vs-cross-validation-in-machine-learning-7637112d3f8f>)

<https://machinelearningmastery.com/machine-learning-ensembles-with-r/> (<https://machinelearningmastery.com/machine-learning-ensembles-with-r/>)

<https://rpubs.com/njvijay/16444> (<https://rpubs.com/njvijay/16444>)

<https://topepo.github.io/caret/model-training-and-tuning.html> (<https://topepo.github.io/caret/model-training-and-tuning.html>)

<https://www.wright.edu/center-for-teaching-and-learning/blog/article/creating-easy-narrated-screen-recordings-on-almost-any-mac> (<https://www.wright.edu/center-for-teaching-and-learning/blog/article/creating-easy-narrated-screen-recordings-on-almost-any-mac>)

<https://cran.r-project.org/web/packages/caretEnsemble/vignettes/caretEnsemble-intro.html> (<https://cran.r-project.org/web/packages/caretEnsemble/vignettes/caretEnsemble-intro.html>)

<https://blog.revolutionanalytics.com/2015/10/the-5th-tribe-support-vector-machines-and-caret.html> (<https://blog.revolutionanalytics.com/2015/10/the-5th-tribe-support-vector-machines-and-caret.html>)

https://uc-r.github.io/naive_bayes (https://uc-r.github.io/naive_bayes)

<https://machinelearningmastery.com/machine-learning-ensembles-with-r/> (<https://machinelearningmastery.com/machine-learning-ensembles-with-r/>)

<https://topepo.github.io/caret/available-models.html> (<https://topepo.github.io/caret/available-models.html>)

<https://rpubs.com/zxs107020/370699> (<https://rpubs.com/zxs107020/370699>) CaretList and CaretStack

<http://danlec.com/st4k#questions/49725934> (<http://danlec.com/st4k#questions/49725934>)

https://github.com/suesui117/da5030_final_project/blob/main/DA5030.Proj.Sui.html#questions/49725934

<https://towardsdatascience.com/a-comprehensive-machine-learning-workflow-with-multiple-modelling-using-caret-and-caretensemble-in-fcbf6d80b5f2> (<https://towardsdatascience.com/a-comprehensive-machine-learning-workflow-with-multiple-modelling-using-caret-and-caretensemble-in-fcbf6d80b5f2>)

<https://www.neuraldesigner.com/learning/examples/bank-churn> (<https://www.neuraldesigner.com/learning/examples/bank-churn>)

<https://datascience.stackexchange.com/questions/17146/does-ensemble-bagging-boosting-stacking-etc-always-at-least-increase-perfor> (<https://datascience.stackexchange.com/questions/17146/does-ensemble-bagging-boosting-stacking-etc-always-at-least-increase-perfor>)

https://stats.libretexts.org/Bookshelves/Computing_and_Modeling/RTG%3A_Classification_Methods/4%3A_Numerical_Experiments_and_Real_Data_Analysis/Pr (https://stats.libretexts.org/Bookshelves/Computing_and_Modeling/RTG%3A_Classification_Methods/4%3A_Numerical_Experiments_and_Real_Data_Analysis/Pr)

https://subscription.packtpub.com/book/big_data_and_business_intelligence/9781788397872/1/ch01lv1sec27/pros-and-cons-of-neural-networks (https://subscription.packtpub.com/book/big_data_and_business_intelligence/9781788397872/1/ch01lv1sec27/pros-and-cons-of-neural-networks)

<https://medium.com/analytics-vidhya/accuracy-vs-f1-score-6258237beca2> (<https://medium.com/analytics-vidhya/accuracy-vs-f1-score-6258237beca2>)

<https://www.machinelearningplus.com/machine-learning/caret-package/> (<https://www.machinelearningplus.com/machine-learning/caret-package/>) (one hot encoding, caretStack)

https://www.saedsayad.com/decision_tree_reg.htm (https://www.saedsayad.com/decision_tree_reg.htm) (Decision Tree)

<https://machinelearningmastery.com/why-one-hot-encode-data-in-machine-learning/> (<https://machinelearningmastery.com/why-one-hot-encode-data-in-machine-learning/>) (Integer encoding vs one hot encoding)